

Introduction to MCMC

Bayesian Psychometric Models, Lecture 7

The current lecture is built loosely off of Levy and Mislevy (2016) Chapter 5, with information from Chapters 4 and 6 (although not necessarily all from the book).

Introduction to Markov Chain Monte Carlo Estimation

Bayesian analysis is all about estimating the posterior distribution

- Up until now, we have worked with posterior distributions that fairly well-known
 - Beta-Binomial had a Beta distribution
 - In general, likelihood distributions from the exponential family have conjugate priors
 - * Conjugate prior: the family of the prior is equivalent to the family of the posterior
- Most of the time, however, posterior distributions are not easily obtainable
 - No longer able to use properties of the distribution to estimate parameters
- It is possible to use an optimization algorithm (e.g., Newton-Raphson or Expectation-Maximization) to find maximum value of posterior distribution
 - But, such algorithms may take a very long time for high-dimensional problems
- Instead: “sketch” the posterior by sampling from it – then use that sketch to make inferences
 - Sampling is done via MCMC

Markov Chain Monte Carlo

- MCMC algorithms iteratively sample from the posterior distribution
 - For fairly simplistic models, each iteration has independent samples
 - Most models have some layers of dependency included
 - * Can slow down sampling from the posterior
- There are numerous variations of MCMC algorithms
 - Most of these specific algorithms use one of two types of sampling:
 1. Direct sampling from the posterior distribution (i.e. Gibbs sampling)
 - * Often used when conjugate priors are specified
 2. Indirect (rejection-based) sampling from the posterior distribution (e.g., Metropolis-Hastings)
- Efficiency is the main reason for so many algorithms
 - Efficiency in this context: How quickly the algorithm converges and provides adequate coverage (“sketching”) of the posterior distribution
 - No one algorithm is uniformly most efficient for all models (here model = likelihood \times prior)
- The good news is that many software packages (JAGS and MPlus, especially) don’t make you choose which specific algorithm to use
- The bad news is that sometimes your model may take a large amount of time to reach convergence (think days or weeks)
- You can also code your own custom algorithm to make things run smoother

Commonalities Across MCMC Algorithms

- Despite having fairly broad differences regarding how algorithms sample from the posterior distribution, there are quite a few things that are similar across algorithms:
 1. A period of the Markov chain where sampling is not directly from the posterior
 - The burnin period

- 2. Methods used to assess convergence of the chain to the posterior distribution
 - Often involving the need to use multiple chains with independent and differing starting values
- 3. Summaries of the posterior distribution
- Further, rejection-based sampling algorithms often need a tuning period to make the sampling more efficient
 - The tuning period comes before the algorithm begins its burnin period

To demonstrate each type of algorithm, we will use a model for a normal distribution (Chapter 4 of Levy & Mislevy, 2016) - We will investigate each, briefly - We will then switch over to JAGS to show the syntax and let JAGS work

We will conclude by talking about assessing convergence and how to report parameter estimates.

Example Data: Post-Diet Weights

Borrowing an example from (<https://stats.idre.ucla.edu/spss/library/spss-libraryhow-do-i-handle-interactions-of-continuous-andcategorical-variables/>), the file DietData.csv contains data from 30 respondents who participated in a study regarding the effectiveness of three types of diets. Variables in the data set are:

1. Respondent: Respondent number 1-30
2. DietGroup: A 1, 2, or 3 representing the group to which a respondent was assigned
3. HeightIN: The respondent's height in inches
4. WeightLB: The respondent's weight, in pounds, recorded following the study

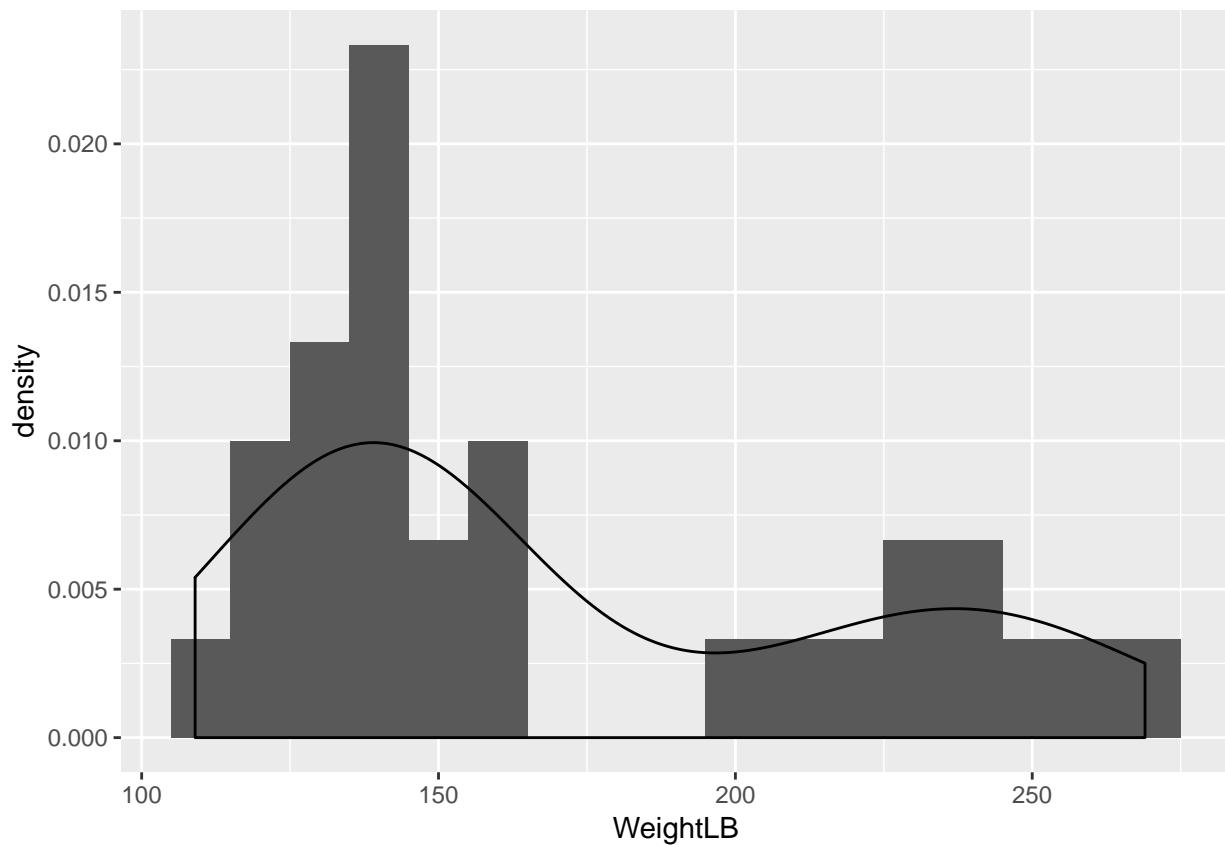
Here, weight is the dependent variable. The motivating research question is:

Are there differences in final weights between the three diet groups, and, if so, what are the nature of the differences?

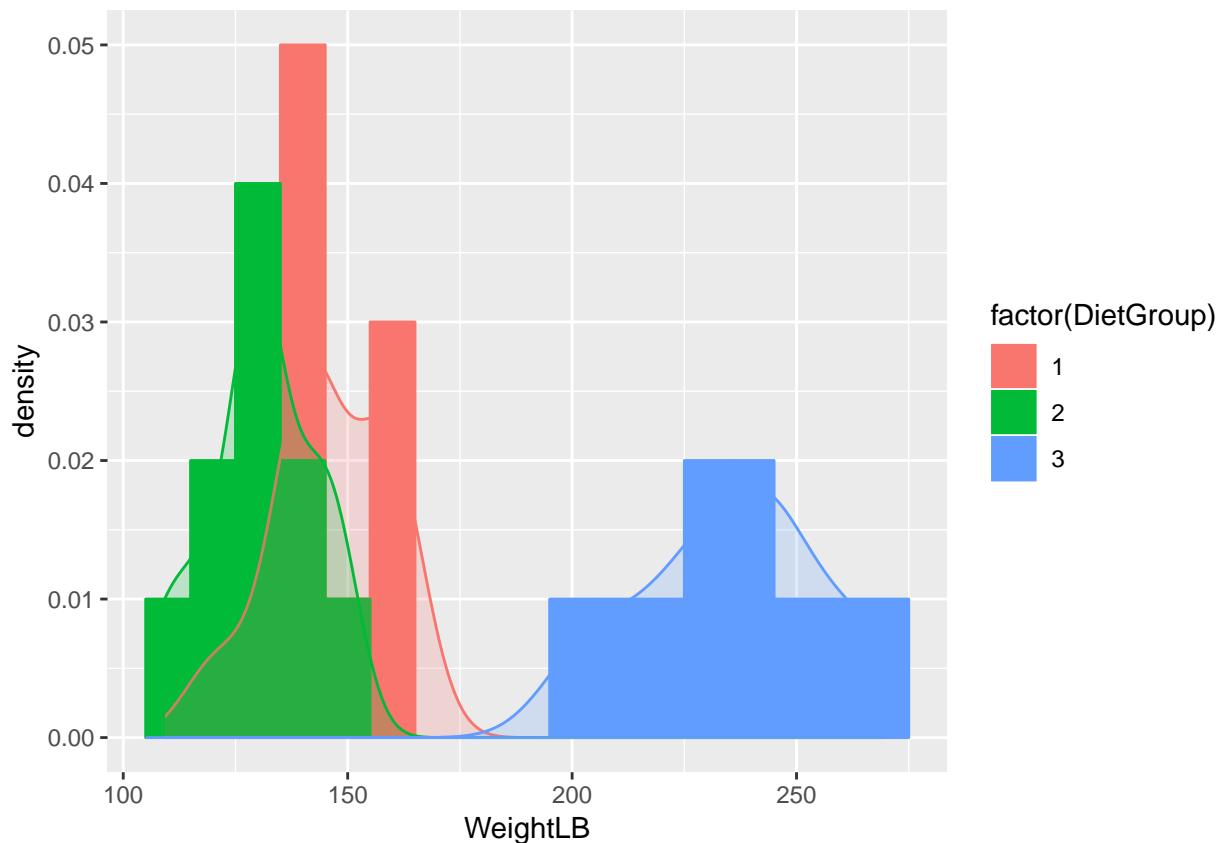
The following syntax reads the data into R and provides some pictures of the data:

```
DietData = read.csv(file = "DietData.csv")

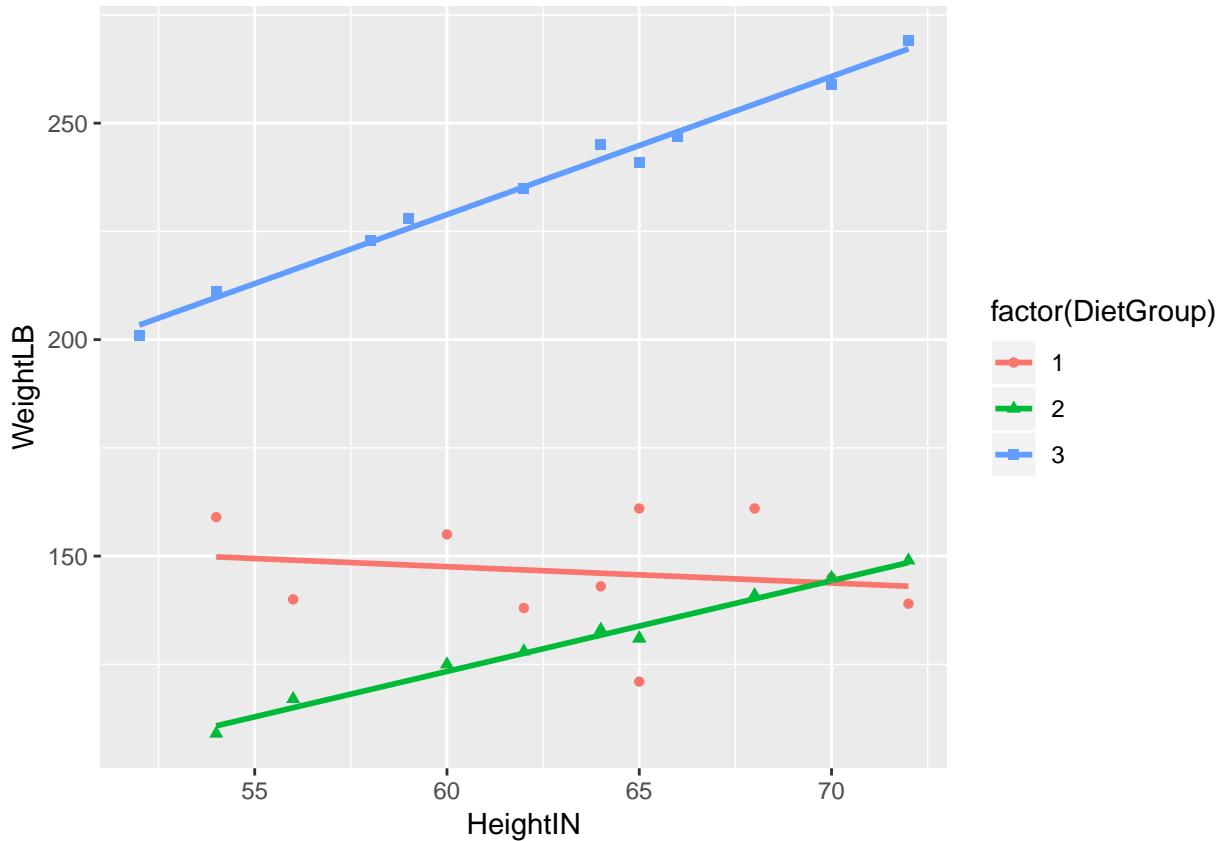
ggplot(data = DietData, aes(x = WeightLB)) +
  geom_histogram(aes(y = ..density..), position = "identity", binwidth = 10) +
  geom_density(alpha=.2)
```



```
ggplot(data = DietData, aes(x = WeightLB, color = factor(DietGroup), fill = factor(DietGroup))) +  
  geom_histogram(aes(y = ..density..), position = "identity", binwidth = 10) +  
  geom_density(alpha=.2)
```



```
ggplot(data = DietData, aes(x = HeightIN, y = WeightLB, shape = factor(DietGroup), color = factor(DietG
```



Now, your turn to answer questions:

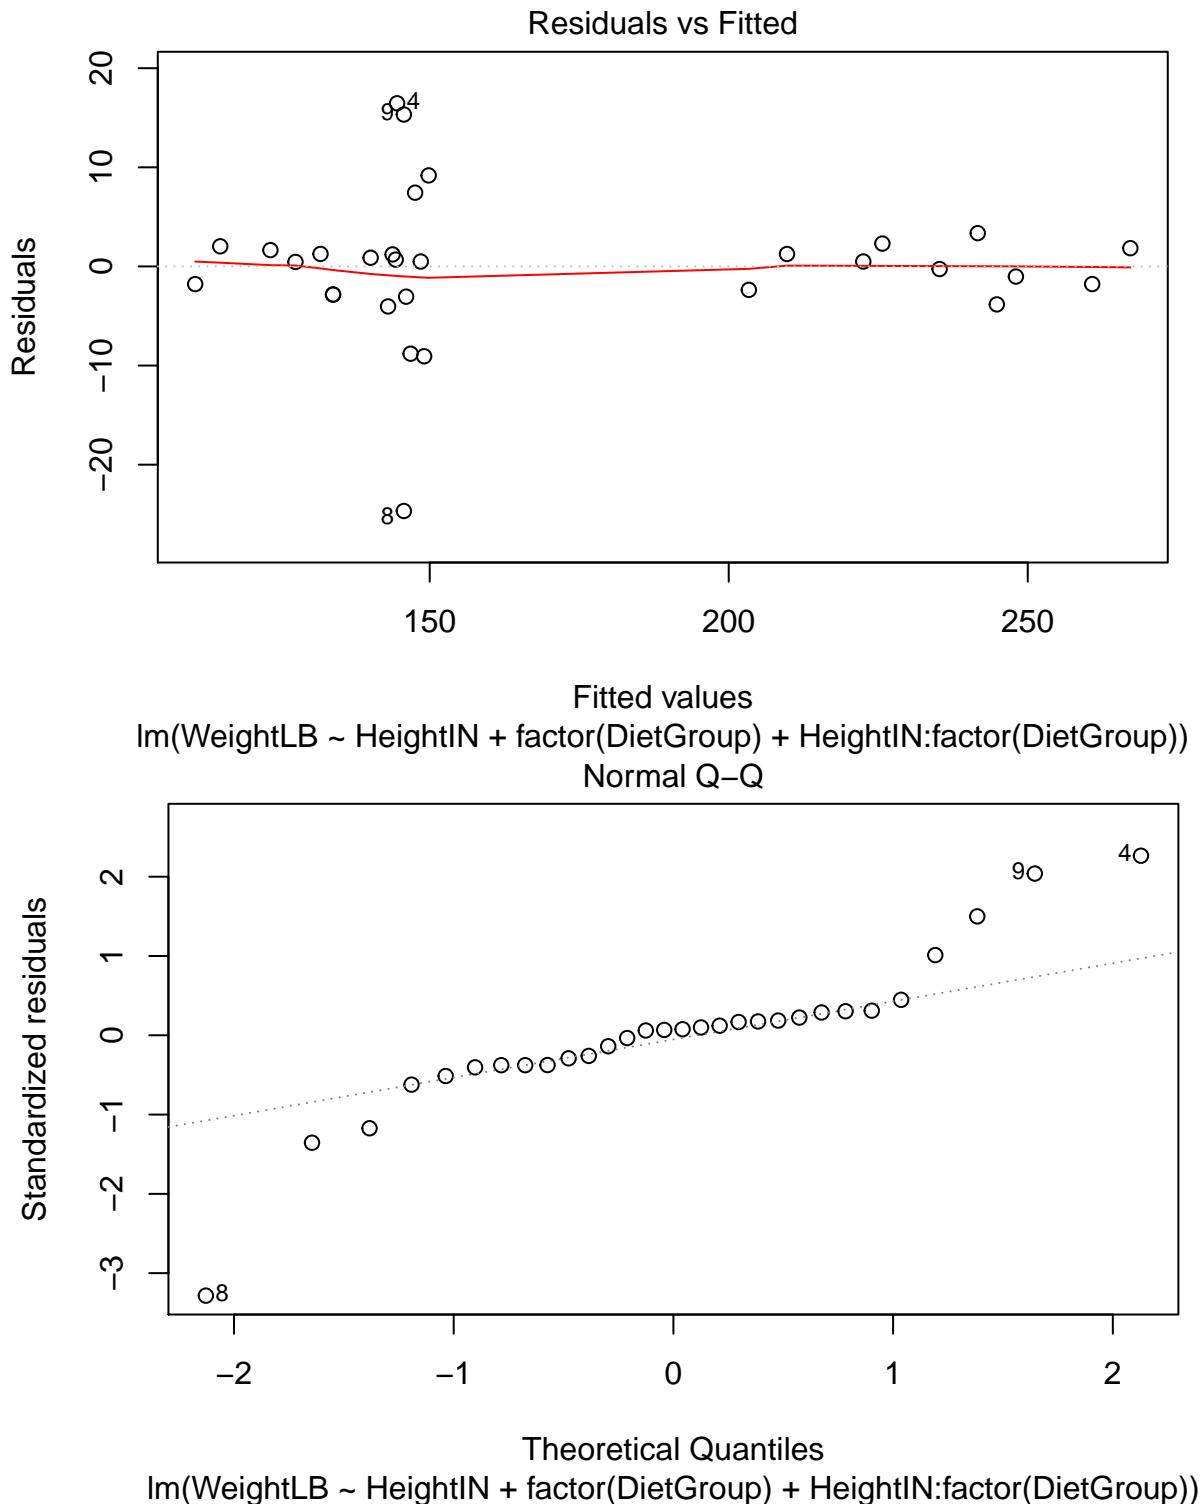
1. What type of analysis seems most appropriate for these data?
2. Is the dependent variable (WeightLB) is appropriate as-is for such analysis or does it need transformed?

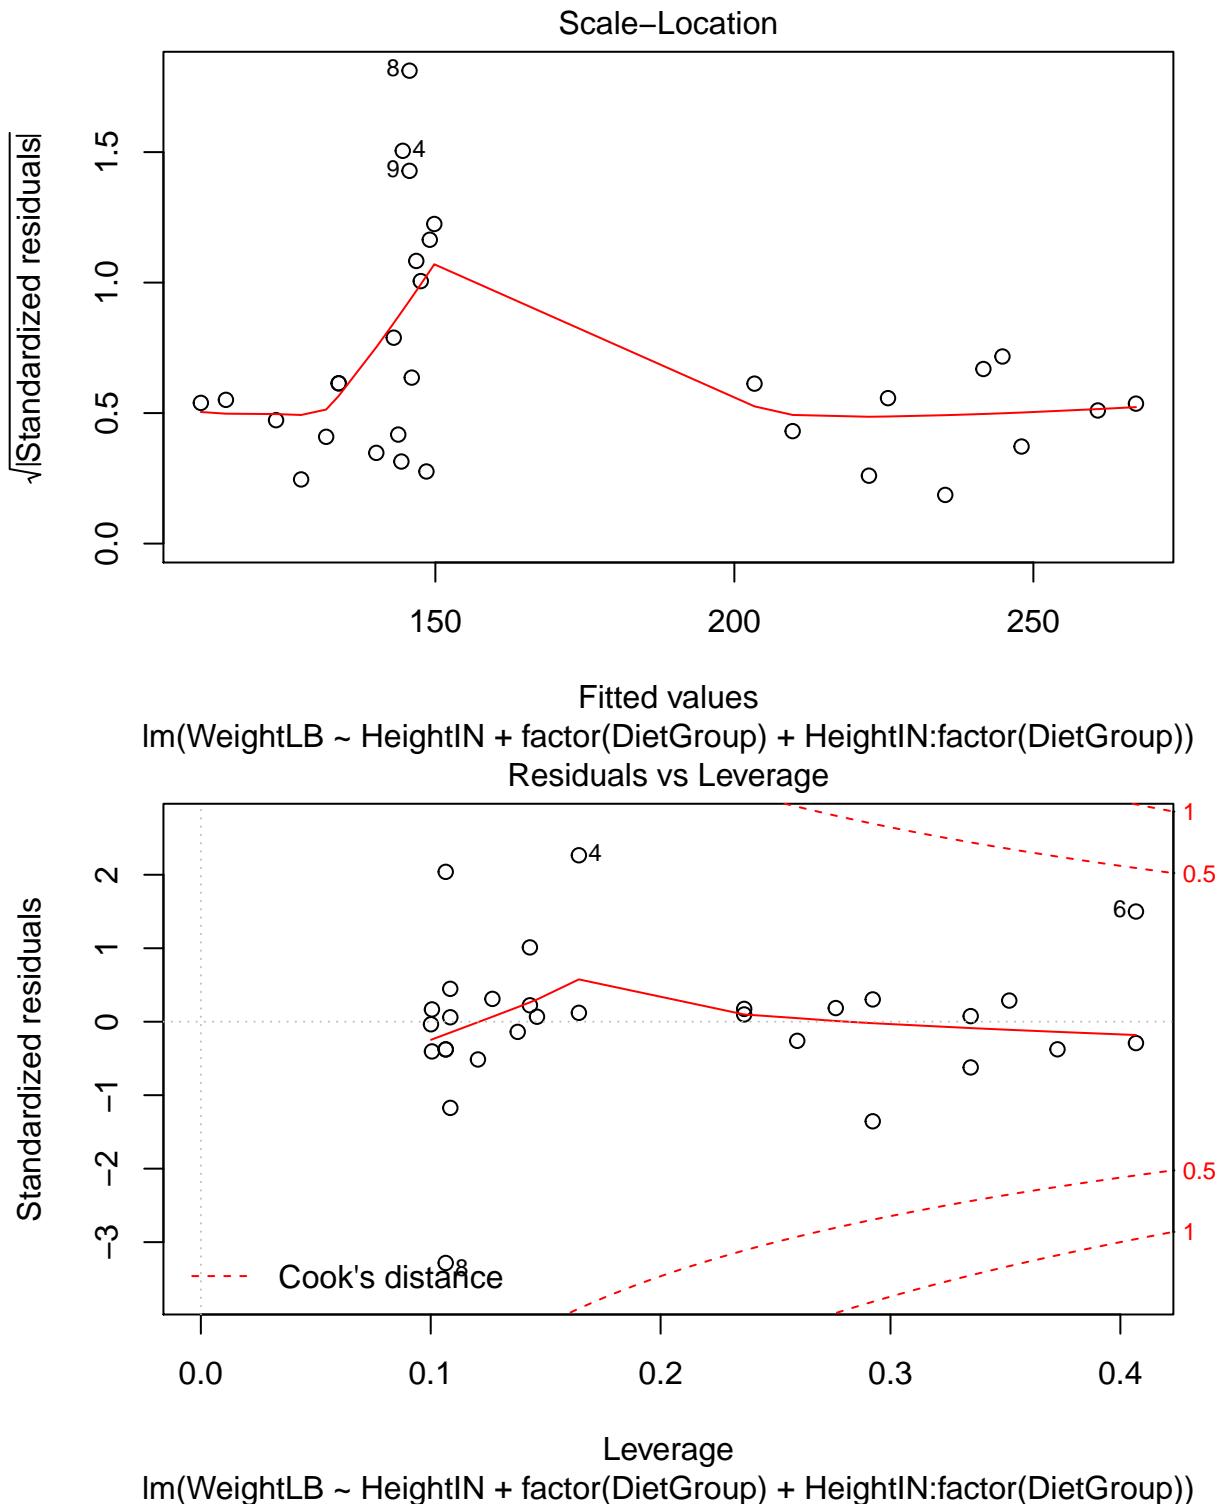
```
# full analysis model suggested by data:
```

```
FullModel = lm(formula = WeightLB ~ HeightIN + factor(DietGroup) + HeightIN:factor(DietGroup), data = D)
```

```
# examining assumptions and leverage of fit
```

```
plot(FullModel)
```





```
# looking at ANOVA table
anova(FullModel)
```

```
## Analysis of Variance Table
##
## Response: WeightLB
##
```

Df	Sum Sq	Mean Sq	F value	Pr(>F)
----	--------	---------	---------	--------

```

## HeightIN             1     684      684  10.819  0.003092 ***
## factor(DietGroup)    2   66726    33363 527.950 < 2.2e-16 ***
## HeightIN:factor(DietGroup) 2    2186     1093  17.293 2.234e-05 ***
## Residuals            24    1517      63
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
# looking at parameter summary
summary(FullModel)

##
## Call:
## lm(formula = WeightLB ~ HeightIN + factor(DietGroup) + HeightIN:factor(DietGroup),
##      data = DietData)
##
## Residuals:
##       Min     1Q Median     3Q    Max
## -24.6724 -2.7169  0.4958  1.7918 16.4581
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)                170.1664   29.2786   5.812 5.42e-06 ***
## HeightIN                  -0.3768    0.4587  -0.822 0.419392
## factor(DietGroup)2        -172.5639   41.4062  -4.168 0.000345 ***
## factor(DietGroup)3        -132.6675   38.7846  -3.421 0.002241 **
## HeightIN:factor(DietGroup)2  2.4727    0.6486   3.812 0.000846 ***
## HeightIN:factor(DietGroup)3  3.5666    0.6132   5.817 5.36e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.949 on 24 degrees of freedom
## Multiple R-squared:  0.9787, Adjusted R-squared:  0.9742
## F-statistic: 220.3 on 5 and 24 DF,  p-value: < 2.2e-16

```

Bayesian Analysis of Our Example Data

Let's examine how our example data should be analyzed using Bayesian methods. To do so we need to specify:

1. Likelihood function (from the model)
2. Prior distributions for all model parameters

Linear Models Likelihood Function (Scalar Version)

The classic linear model is one where, for a respondent r ($r = 1, \dots, N$), a dependent variable (or outcome), y_r , is predicted via a set of independent variables (or predictors), x_{rv} where $v = 1, \dots, V$ and, sometimes, their interactive product, by means of a set of regression coefficients β_v :

$$Y_r = \beta_0 + \beta_1 x_{r1} + \beta_2 x_{r2} + \cdots + \beta_V x_{rV} + e_r = \mathbf{x}_r \boldsymbol{\beta} + e_r$$

Here, e_r is the residual or error term for respondent r , with $e_r \sim N(0, \sigma_e^2)$.

So, loosely, we can say that an observation $Y_r \sim N(\mathbf{x}_r \boldsymbol{\beta}, \sigma_e^2)$. This now gives us the set of parameters for which we need to specify prior distributions: $\boldsymbol{\beta}$ and σ_e^2 .

Linear Models: Conjugate Prior Distributions

Conjugate priors are prior distributions that have the same form as the posterior distribution. As we will see in another lecture, conjugate priors are helpful because they make our Bayesian sampling algorithms run efficiently. We will stay with conjugate prior distributions for this lecture.

Choosing conjugate prior distributions in Bayesian linear models has one complicating factor: The likelihood for β is conditional on σ_e^2 . So, let's start there. In particular:

$$f(\beta, \sigma_e^2) = f(\beta | \sigma_e^2) f(\sigma_e^2)$$

Prior for σ_e^2

To make sense of this, we can start with σ_e^2 , which has a conjugate prior distribution that is inverse-gamma. The inverse-gamma distribution (see (https://en.wikipedia.org/wiki/Inverse-gamma_distribution)[https://en.wikipedia.org/wiki/Inverse-gamma_distribution]) is one that has support of $(0, \infty]$, which matches the range of σ_e^2 . The distribution has two parameters, α (sometimes called the shape; $\alpha > 0$) and β (sometimes called the scale, although in the `invgamma` R package, it is the rate, $\beta > 0$).

To get a sense of how these parameters function, the mode is a good measure of central tendency, which is $\frac{\beta}{\alpha+1}$. In practice, we can rephrase these in terms of our expected degrees of freedom total ν_0 and our expected residual variance σ_{e0}^2 :

$$\alpha_0 = \frac{\nu_0}{2}$$

$$\beta_0 = \frac{\nu_0 \sigma_0^2}{2}$$

In the plot below, I use the values from our classical linear model analysis to provide an example of what the prior looks like:

```
nu.0 = summary(FullModel)$df[2]
sigma2.0 = summary(FullModel)$sigma^2

alpha.0 = nu.0/2
beta.0 = nu.0*sigma2.0/2

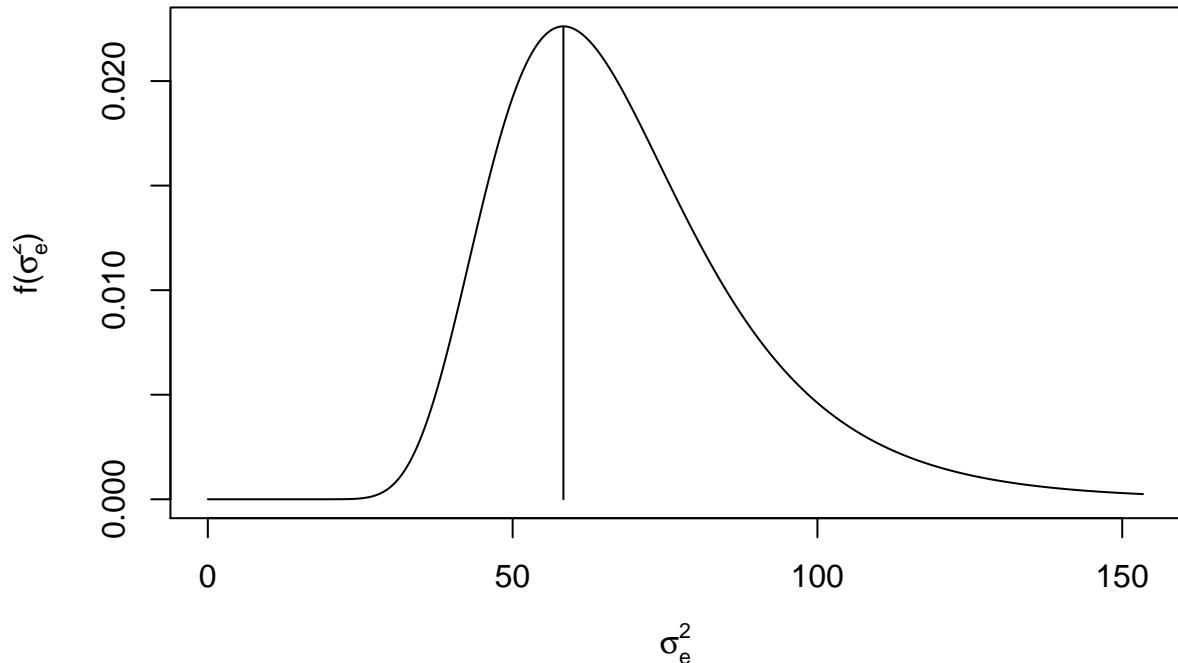
PriorMode = beta.0/(alpha.0+1)

maxX = qinvgamma(.995, shape = alpha.0, rate= beta.0)
minX = 0.001
step = (maxX - minX)/10000

sigma2 = seq(minX, maxX, step)
height = dinvgamma(x = sigma2, shape = alpha.0, rate = beta.0)

plot(x = sigma2, y = height, type = "l", ylab = expression(paste("f(", sigma[e]^2, ")")), xlab = expression(
    main = paste0("Prior Mode:", round(PriorMode, 2), "; MLE:", round(sigma2.0, 2)))
lines(x = c(PriorMode, PriorMode), y = c(0, dinvgamma(x = PriorMode, shape = alpha.0, rate = beta.0)))
```

Prior Mode:58.33; MLE:63.19



One complicating factor is that for normal distributions (where we will need to put the parameter σ_e^2 with this prior) JAGS uses a “precision” parameter of $\frac{1}{\sigma_e^2}$ rather than a “dispersion” of σ_e^2 . Therefore, we must specify our prior for $\tau_e = \frac{1}{\sigma_e^2}$, which follows a gamma distribution.

The good news is that the parameters of the gamma distribution are the same, however, you have to check what they are called in R. Here is how the previous plot looks with a gamma distribution:

```

nu.0 = summary(FullModel)$df[2]
sigma2.0 = summary(FullModel)$sigma^2
tau.e.0 = 1/sigma2.0

alpha.0 = nu.0/2
beta.0 = nu.0*sigma2.0/2

PriorMode = (alpha.0-1)*(1/beta.0)

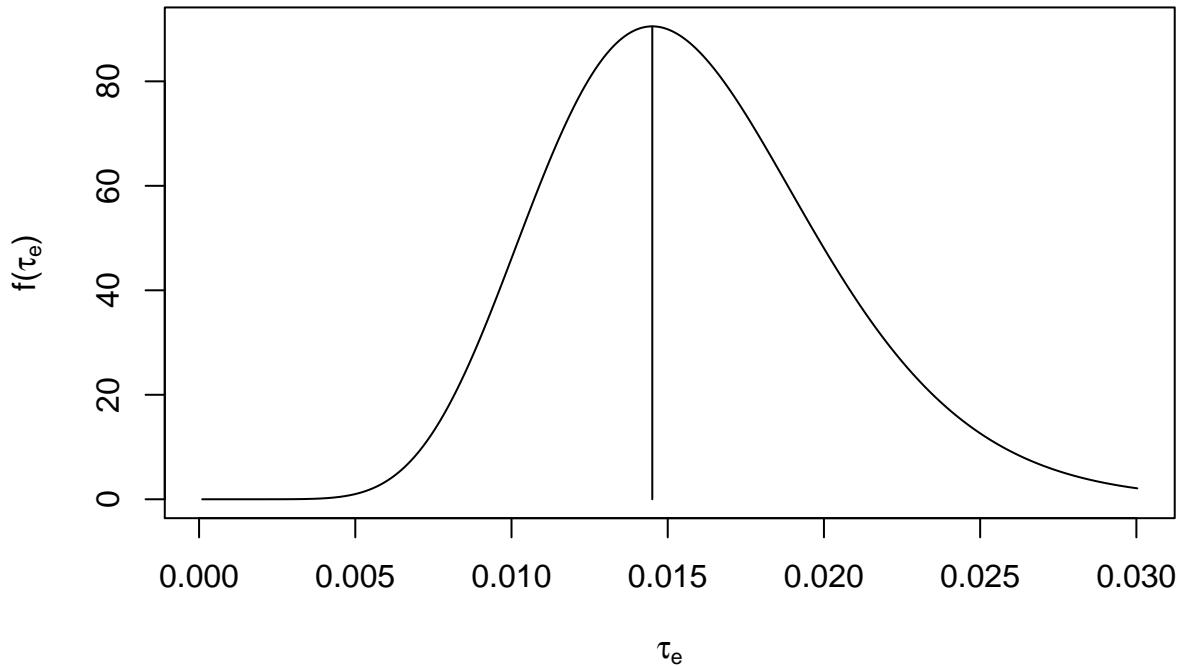
maxX = qgamma(.995, shape = alpha.0, rate= beta.0)
minX = 0.0001
step = .00001

tau.e = seq(minX, maxX, step)
height = dgamma(x = tau.e, shape = alpha.0, rate = beta.0)

plot(x = tau.e, y = height, type = "l", ylab = expression(paste("f(", tau[e], ")")), xlab = expression(
  main = paste0("Prior Mode:", round(PriorMode, 2), "; MLE:", round(1/sigma2.0, 2)))
lines(x = c(PriorMode, PriorMode), y = c(0, dgamma(x = PriorMode, shape = alpha.0, rate = beta.0)))

```

Prior Mode:0.01; MLE:0.02



“Uninformative” Choices of Conjugate Prior for σ_e^2 (really, τ_e)

One way to make an “uninformative” but conjugate prior for σ_e^2 is to recall the terms we used to specify the prior parameters: ν_0 as the expected DF total and σ_{e0}^2 as the expected residual variance.

For the residual variance: 1. The maximum residual variance is the marginal variance of y : occurs when $\beta = \mathbf{0}$ 2. Higher values of variance parameters result in higher posterior standard deviations for β , meaning they are more conservative

So, we’ll start with an estimate of $\sigma_{e0}^2 = \sigma_y^2$. For our sample, $\sigma_y^2 = 2452.14$.

Next, recall that ν_0 is the expected degrees of freedom total for the analysis. As degrees of freedom decrease, the asymptotic sampling error of σ_e^2 gets large, so you can pick a value of ν_0 that is small (but above 0). See the following plot (for σ_e^2 , but the parameter choices apply to τ):

```
sigma2.0.U = var(DietData$WeightLB)
nu.0.U = 1

alpha.0.U = nu.0.U/2
beta.0.U = (nu.0.U*sigma2.0.U)/2

nu.0 = summary(FullModel)$df[2]
sigma2.0 = summary(FullModel)$sigma^2

alpha.0 = nu.0/2
beta.0 = (nu.0*sigma2.0)/2

PriorMode = beta.0/(alpha.0+1)

maxX = 3*sigma2.0.U
minX = 0.001
```

```

step = (maxX - minX)/10000

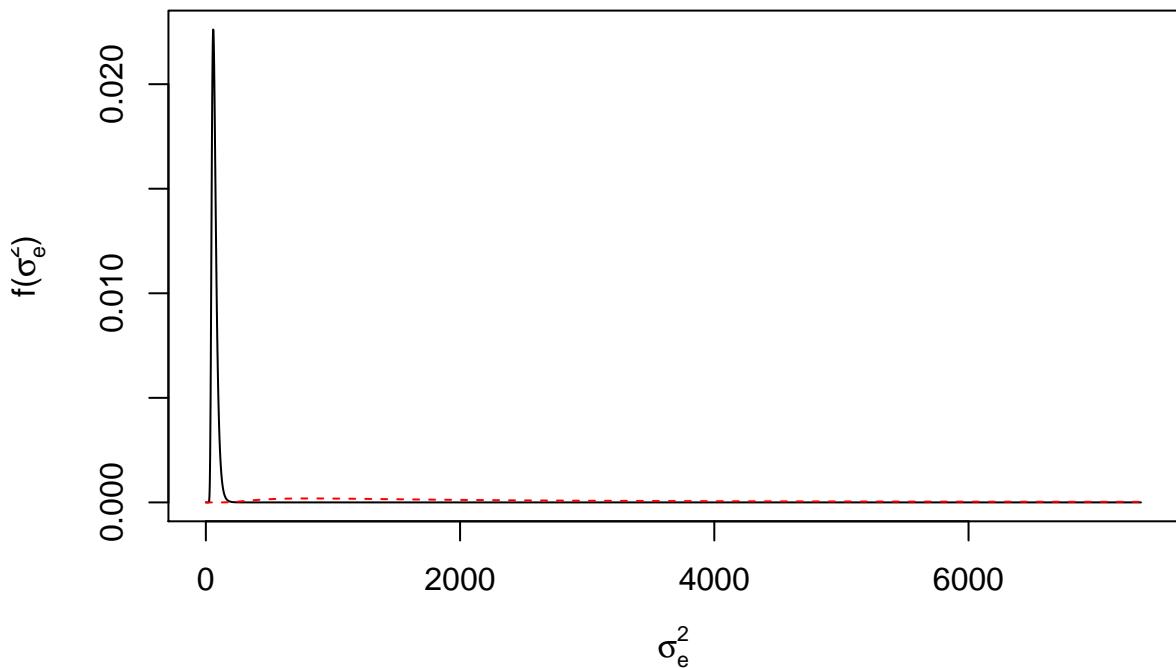
sigma2 = seq(minX, maxX, step)

height = dinvgamma(x = sigma2, shape = alpha.0, rate = beta.0)
heightU = dinvgamma(x = sigma2, shape = alpha.0.U, rate = beta.0.U)

matplot(
  x = cbind(sigma2, sigma2),
  y = cbind(height, heightU),
  type = "l",
  ylab = expression(paste("f(", sigma[e] ^ 2, ")")),
  xlab = expression(sigma[e] ^ 2),
  main = "Red is uninformative prior; black is prior based on suggestion in book"
)

```

Red is uninformative prior; black is prior based on suggestion in book

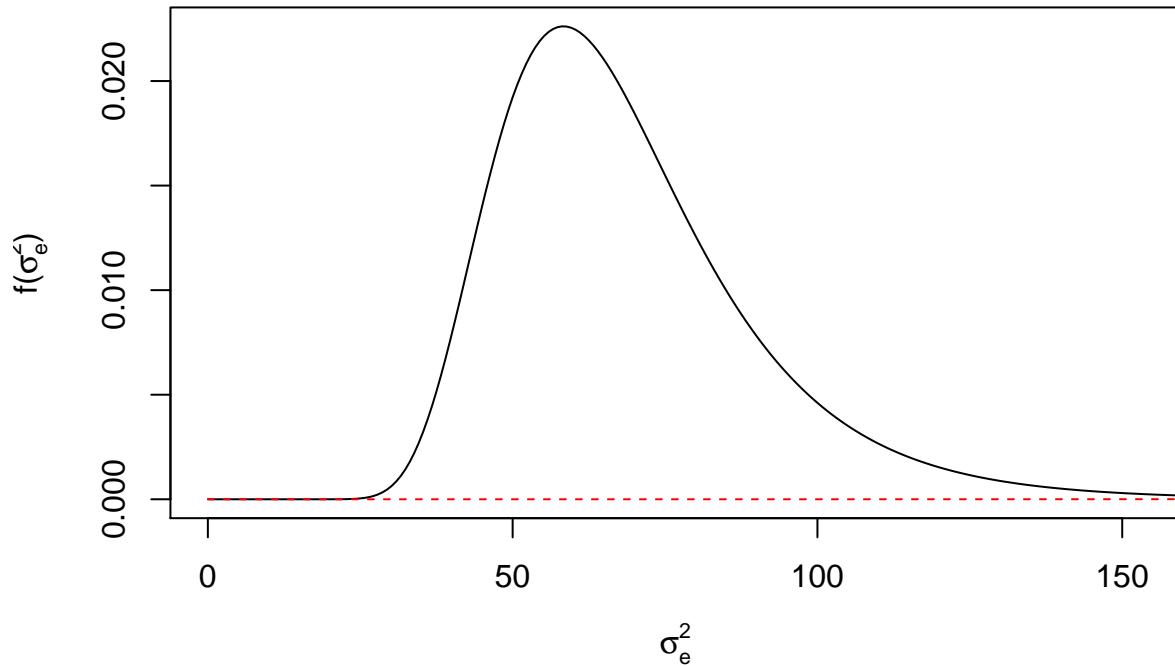


```

matplot(
  x = cbind(sigma2, sigma2),
  y = cbind(height, heightU),
  type = "l",
  ylab = expression(paste("f(", sigma[e] ^ 2, ")")),
  xlab = expression(sigma[e] ^ 2),
  main = "Same Function: Close Up Near Original Prior",
  xlim = c(0, qinvgamma(.995, shape = alpha.0, rate = beta.0))
)

```

Same Function: Close Up Near Original Prior



Prior for β

Conditional on a value of σ_e^2 , the conjugate prior for β follows a multivariate normal distribution: $\beta \sim MVN(\bar{\beta}, \Sigma_\beta)$. Often, this is expressed as a set of independent univariate priors where the mean for each β is given along with some type of variance. That said, the use of normal priors in JAGS uses precision instead of variance (so $1/\text{variance}$).

These priors can be made to be uninformative by selecting very low values for precision (meaning very high levels of variance). Remember, 99% of the mass of a univariate normal distribution falls within $\pm 3\text{SD}$.

That said, the scale of the regression coefficients will not be constant, so if picking one value for a hyperparameter for the prior of each, pick a very large variance

Bayesian Linear Models Analysis with JAGS

The first step to building our Bayesian linear model analysis in JAGS is to put all of our data and prior values into an R list object. Here, we will

```
# specify specs for analysis
n.chains = 4

# specify values of hyperparameters for prior distributions

# for sigma2_e (tau)
sigma2.0 = var(DietData$WeightLB)
nu.0 = 1

alpha.tau.0 = nu.0/2
beta.tau.0 = (nu.0*sigma2.0)/2
```

```

# for betas:
betaMean.0 = 0
betaVariance.0 = 100000000
betaTau.0 = 1/betaVariance.0

# take data frame and create model matrix -- which gives us the columns of X:
BayesianFullModel.Matrix = model.matrix(object = FullModel)

# append with weight and then select only columns needed for model:
BayesianFullModel.Matrix = data.frame(cbind(DietData$WeightLB, BayesianFullModel.Matrix))
names(BayesianFullModel.Matrix) = c("weight", "intercept", "height", "group2", "group3", "hg2", "hg3")
BayesianFullModel.Matrix = BayesianFullModel.Matrix[c("weight", "height", "group2", "group3")]

# add data and N to list object that we will pass to JAGS
BayesianFullModel.JAGS.Data = list(N = nrow(DietData),
                                    weight = BayesianFullModel.Matrix$weight,
                                    height = BayesianFullModel.Matrix$height,
                                    group2 = BayesianFullModel.Matrix$group2,
                                    group3 = BayesianFullModel.Matrix$group3,
                                    betaMean.0 = betaMean.0,
                                    betaTau.0 = betaTau.0,
                                    alpha.tau.0 = alpha.tau.0,
                                    beta.tau.0 = beta.tau.0)

```

Next, we can create our JAGS model syntax. We will use the `R2jags` package, so we enclose the syntax in a function. Here, we put the names of the variables and hyperparameters we created in our previous step directly into the JAGS model syntax.

There are several things to note here:

- Hyperparameter values are explicit numbers in syntax, not variables entered into JAGS
- The example syntax on p. 129 has a calculation of R^2 – however, it uses the MLE for σ_y^2 and a version of the MLE for σ_e^2 , so it is not included

`BayesianFullModel.Syntax = "`

```

model{

  # prior distributions:
  # Note that terms on the left are model parameter names we create here
  beta.0 ~ dnorm(betaMean.0, betaTau.0) # prior for the intercept
  beta.height ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional slope of height
  beta.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 2 and group 1
  beta.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 3 and group 1
  beta.height.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 2
  beta.height.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 3
  tau.e ~ dgamma(alpha.tau.0, beta.tau.0)      #prior for 1/sigma^2_e, residual precision

  # conditional distribution of the data (uses priors above)
  for (i in 1:N){

    # creating conditional mean to put into model statement
    weight.hat[i] <- beta.0 + beta.height*height[i] + beta.group2*group2[i] + beta.group3*group3[i] +

```

```

# error values (for R^2)
error[i] = weight[i]-weight.hat[i]

# likelihood from model:
weight[i] ~ dnorm(weight.hat[i], tau.e)

}

# created values to track throughout the Markov chain
sigma2.e <- 1/tau.e                      #create residual variance from 1/precision
sigma.e   <- sqrt(sigma2.e)                 #create residual SD

# calculate mean difference between groups 2 and 3:
dif.group2.group3 <- beta.group3 - beta.group2

# calculate R^2 based on http://www.stat.columbia.edu/~gelman/research/unpublished/bayes_R2.pdf
variance.pred = sd(weight.hat[])*sd(weight.hat[])
variance.error = sd(error[])*sd(error[])
Rsquare = variance.pred/(variance.pred + variance.error)
}

"

# the parameter names come from the model syntax in the previous step
BayesianFullModel.Parameters = c("beta.0",
                                 "beta.height",
                                 "beta.group2",
                                 "beta.group3",
                                 "beta.height.group2",
                                 "beta.height.group3",
                                 "dif.group2.group3",
                                 "tau.e",
                                 "sigma2.e",
                                 "sigma.e",
                                 "Rsquare",
                                 "deviance")

# in order to have reproducible Markov chains, we need to initialize the random number generator and set
# pick a number: Here is the date I created this syntax
BayesianFullModel.Seed = 06022019

# Note: here, the random number seed cannot be the same per seed or the chains will be the same
RNGname = c("Wichmann-Hill", "Marsaglia-Multicarry", "Super-Duper", "Mersenne-Twister")
if (RNGname[1] %in% c("Wichmann-Hill", "Marsaglia-Multicarry",
                     "Super-Duper", "Mersenne-Twister")) {
  RNGname[1] <- paste("base::", RNGname[1], sep = "")
}

BayesianFullModel.Init.Values <- vector("list", n.chains)
for (i in 1:n.chains) {
  BayesianFullModel.Init.Values[[i]]$.RNG.name <- RNGname[1]
  BayesianFullModel.Init.Values[[i]]$.RNG.seed <- BayesianFullModel.Seed + i
}

```

Next, to run JAGS, we need a few more values to save: (1) the names of the parameters we wish to have

JAGS track across the chains and (2) a random number seed for reproducability. Once we have that, we can run JAGS. Below, I will use the `jags()` function, to which we should also provide the number of chains, number of iterations total, and number of burnin iterations.

```
# load two JAGS modules: GLM (which makes the analysis go more efficiently) and DIC (which gives an ind
load.module("glm")

## module glm loaded

load.module("dic")

## module dic loaded

# Submit model to JAGS for compiling. We'll turn adaptation off to control it in the next line of syntax
BayesianFullModel.JAGS = jags.model(file = textConnection(BayesianFullModel.Syntax), data = BayesianFull

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 30
## Unobserved stochastic nodes: 7
## Total graph size: 259
##
## Initializing model
# adapting the model: No adaptation needed as Gibbs Sampling is being used -- but good to try in case a
adapt(object = BayesianFullModel.JAGS, n.iter = 1000)

## [1] TRUE

# Draw samples; NOTE: BURNIN IS INCLUDED IN SAMPLES
BayesianFullModel.Samples = coda.samples(model = BayesianFullModel.JAGS, variable.names = BayesianFullM

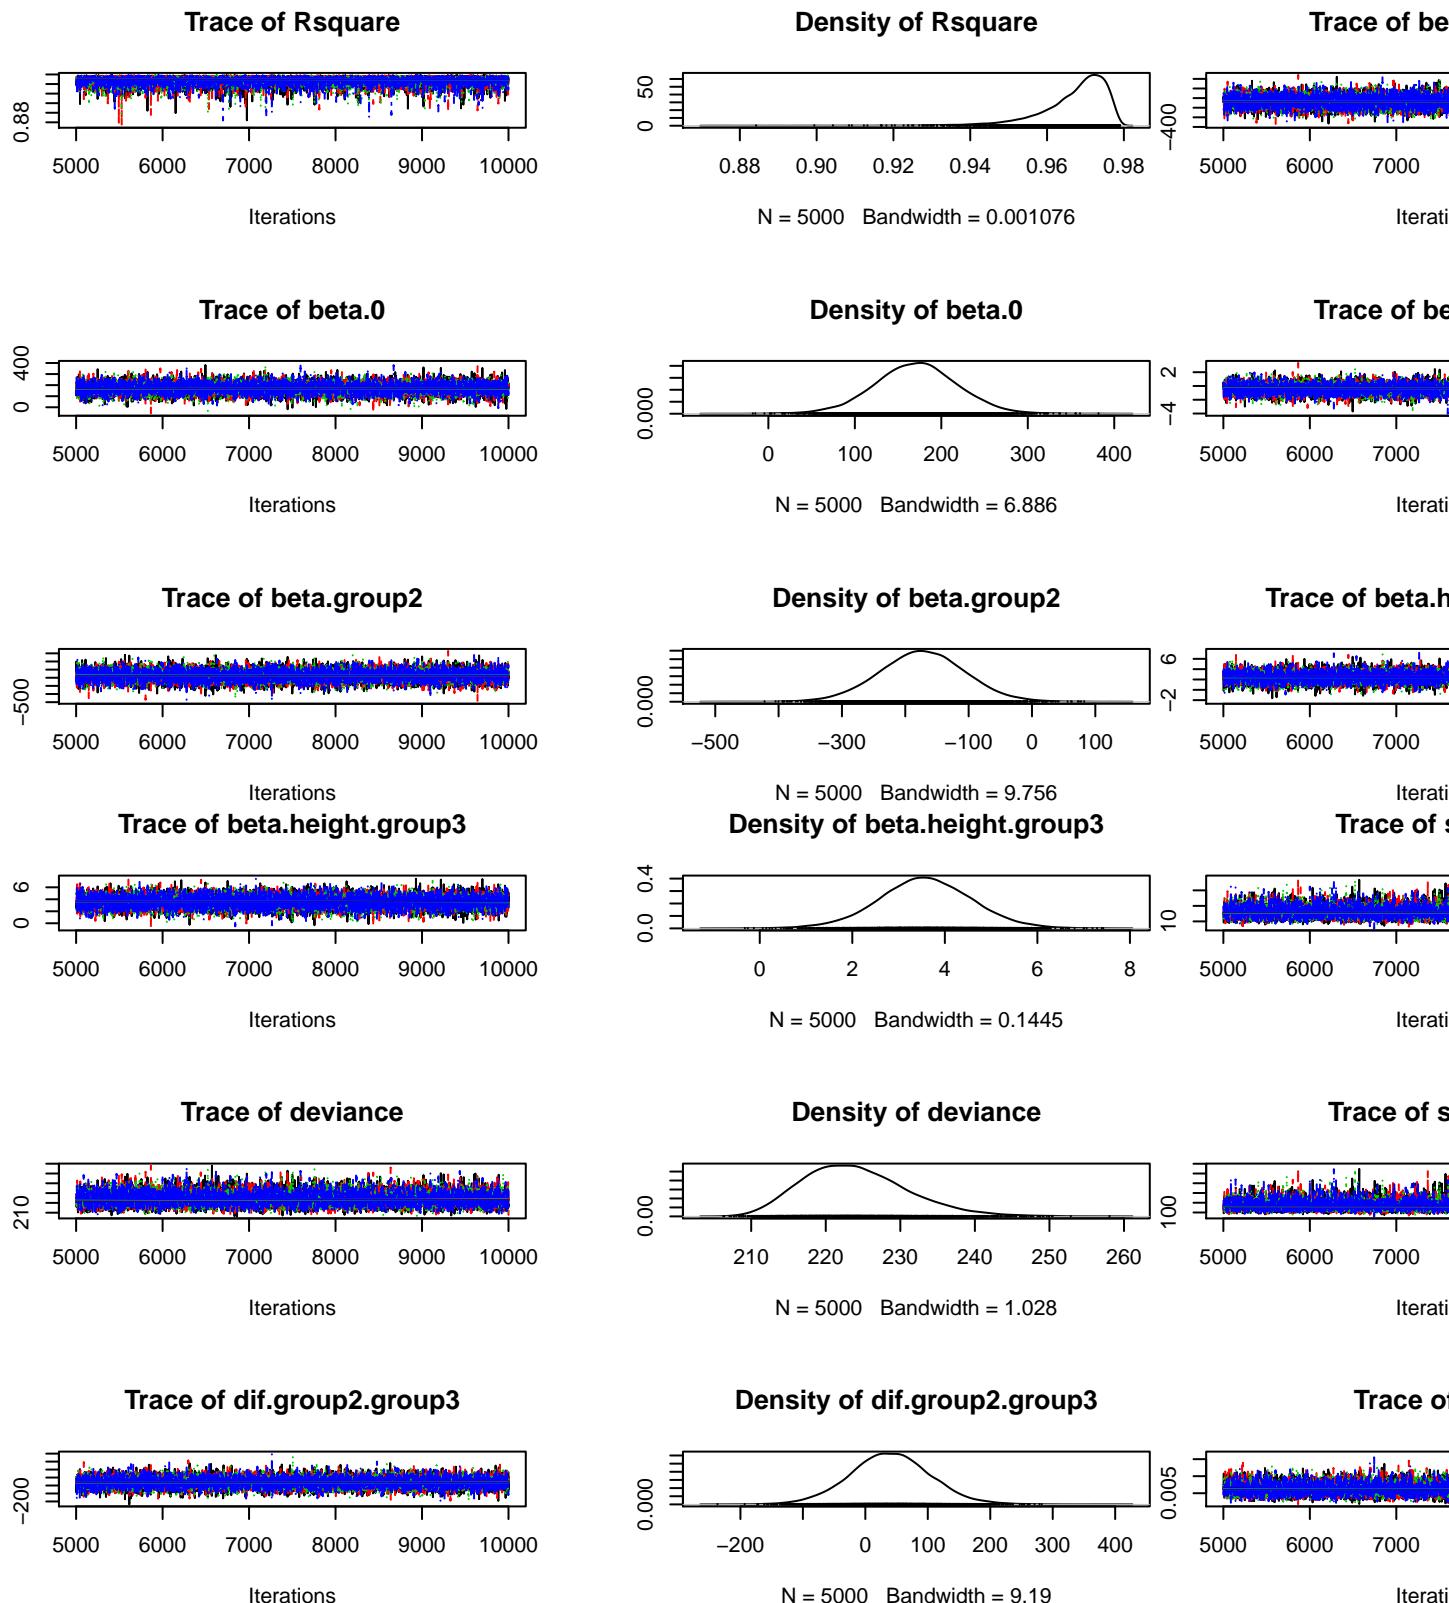
# Remove burnin from samples:
BayesianFullModel.Posterior = window(x = BayesianFullModel.Samples, start = 5001, end = 10000)
```

Checking Chain Convergence

After the model runs, the next step is to assess convergence of the chains to the posterior distribution. This can be accomplished in several ways:

```
# assessing convergence:

# visual assessment
plot(BayesianFullModel.Posterior)
```



```
# using convergence diagnostics:
gelman.diag(BayesianFullModel.Posterior, multivariate = FALSE)
```

```

## Potential scale reduction factors:
##
##          Point est. Upper C.I.
## Rsquare           1       1
## beta.0            1       1
## beta.group2       1       1
## beta.group3       1       1
## beta.height        1       1
## beta.height.group2 1       1
## beta.height.group3 1       1
## deviance          1       1
## dif.group2.group3 1       1
## sigma.e            1       1
## sigma2.e           1       1
## tau.e              1       1

# checking autocorrelation: very low for lag 1
autocorr.diag(BayesianFullModel.Posterior)

##          Rsquare      beta.0   beta.group2 beta.group3   beta.height
## Lag 0    1.000000000 1.000000000 1.000000000 1.000000000 1.000000000
## Lag 1    0.150625821 -0.0008196497 -0.003412160 0.007992317 -0.001429560
## Lag 5    -0.005211171  0.0098458295 -0.001800741 0.008898029  0.007988504
## Lag 10   0.004830756  0.0212639457 -0.000290625 0.004787050  0.021272860
## Lag 50   -0.005125523  0.0010910225  0.012951877 0.001860613  0.002353988
##          beta.height.group2 beta.height.group3   deviance
## Lag 0     1.0000000000          1.0000000000 1.0000000000
## Lag 1    -0.0029447741          0.007692403 0.3234220496
## Lag 5    -0.0024587236          0.007174552 0.0002490533
## Lag 10   -0.0003975335          0.005750982 0.0089525877
## Lag 50   0.0131368090          0.002119942 -0.0044091911
##          dif.group2.group3      sigma.e      sigma2.e      tau.e
## Lag 0     1.0000000000 1.0000000000 1.0000000000 1.0000000000
## Lag 1     0.004809194 0.2023938463 0.2020213503 0.184900242
## Lag 5    -0.008748620 -0.0033246525 -0.0018487243 -0.007869463
## Lag 10   0.002598268 0.0035326988 0.0012221570 0.008525871
## Lag 50   0.003643658 -0.0004026842 0.0003452465 -0.001911239

summary(BayesianFullModel.Posterior)

##
## Iterations = 5001:10000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 5000
##
## 1. Empirical mean and standard deviation for each variable,
## plus standard error of the mean:
##
##          Mean       SD  Naive SE Time-series SE
## Rsquare  9.671e-01 0.009067 6.412e-05   7.474e-05
## beta.0   1.708e+02 48.385970 3.421e-01   3.472e-01
## beta.group2 -1.732e+02 67.877682 4.800e-01   4.723e-01
## beta.group3 -1.333e+02 64.021723 4.527e-01   4.471e-01
## beta.height -3.862e-01 0.757297 5.355e-03   5.431e-03

```

```

## beta.height.group2 2.482e+00 1.062998 7.517e-03      7.469e-03
## beta.height.group3 3.577e+00 1.012075 7.156e-03      6.983e-03
## deviance           2.240e+02 7.032384 4.973e-02      6.654e-02
## dif.group2.group3 3.987e+01 64.521135 4.562e-01      4.605e-01
## sigma.e            1.301e+01 1.932803 1.367e-02      1.679e-02
## sigma2.e           1.730e+02 53.436527 3.779e-01      4.640e-01
## tau.e              6.285e-03 0.001783 1.261e-05      1.512e-05
##
## 2. Quantiles for each variable:
##
##                  2.5%       25%       50%       75%     97.5%
## Rsquare          9.436e-01   0.96342   0.96937   9.733e-01   0.97735
## beta.0           7.393e+01 139.01303 171.02777 2.021e+02 265.44476
## beta.group2     -3.071e+02 -217.85516 -173.00514 -1.285e+02 -40.27762
## beta.group3     -2.596e+02 -175.50529 -132.88512 -9.171e+01 -5.40781
## beta.height      -1.871e+00 -0.87749 -0.38744  1.079e-01  1.12032
## beta.height.group2 3.910e-01  1.78268  2.48167  3.184e+00  4.58555
## beta.height.group3 1.562e+00  2.91743  3.57406  4.242e+00  5.57869
## deviance         2.123e+02 218.93856 223.38443 2.284e+02 239.68384
## dif.group2.group3 -8.691e+01 -2.51107  39.31963  8.169e+01 168.75507
## sigma.e          9.873e+00 11.62916 12.79293 1.413e+01 17.43055
## sigma2.e         9.748e+01 135.23742 163.65905 1.996e+02 303.82395
## tau.e            3.291e-03  0.00501  0.00611  7.394e-03  0.01026

```

The most common method used for assessing convergence is the Gelman-Rubin \hat{R} , which compares the within- and between-chain variance of multiple chains. The `coda` package has a number of convergence diagnostics.

Checking Model-Data Fit: Posterior Predictive Model Checks

A commonly-used method for evaluating the fit of a Bayesian model to the data is to use posterior predictive model checks. Posterior predictive model checking is a process that simulates data using model parameters from the converged section of the Markov chain. The Markov chain is important to use, as some parameters have dependencies in the posterior distribution.

```

nSimData = 1000

# create one big data object to randomly sample from:
BayesianFullModel.Posterior.All = do.call("rbind", BayesianFullModel.Posterior)
parameterNames = colnames(BayesianFullModel.Posterior.All)

# set up model matrix to use in repeated samples
BayesianFullModel.Matrix2 = BayesianFullModel.Matrix[,2:4]

# add columns for intercept and interactions
BayesianFullModel.Matrix2 = as.matrix(data.frame(intercept = matrix(data = 1, nrow = nrow(BayesianFullModel.Matrix2),
height.group2 = BayesianFullModel.Matrix$height*BayesianFullModel.Matrix2,
height.group3 = BayesianFullModel.Matrix$height*BayesianFullModel.Matrix2))
colnames(BayesianFullModel.Matrix2)

## [1] "intercept"      "height"        "group2"        "group3"
## [5] "height.group2"  "height.group3"

parameterNames = c("beta.0", "beta.height", "beta.group2", "beta.group3", "beta.height.group2", "beta.height.group3")

```

```

draws = sample(x = 1:nrow(BayesianFullModel.Posterior.All), size = nSimData, replace = TRUE)

distSummaries = data.frame(mean = matrix(data = NA, nrow = nSimData),
                           variance = matrix(data = NA, nrow = nSimData))

# the following syntax is put into a loop for didactic purposes (a version of apply() would be much faster)
for (i in 1:nSimData){

  # Grab parameters from iteration in draws
  parameters = BayesianFullModel.Posterior.All[draws[i],]

  # Create Beta vector (which becomes mean of random draw), matching parameters with respective columns
  meanVec = BayesianFullModel.Matrix2 %*% matrix(data = parameters[parameterNames], ncol = 1)

  # draw random Y for all observations
  ppcY = rnorm(n = nrow(BayesianFullModel.Matrix2), mean = meanVec, sd = parameters["sigma.e"])

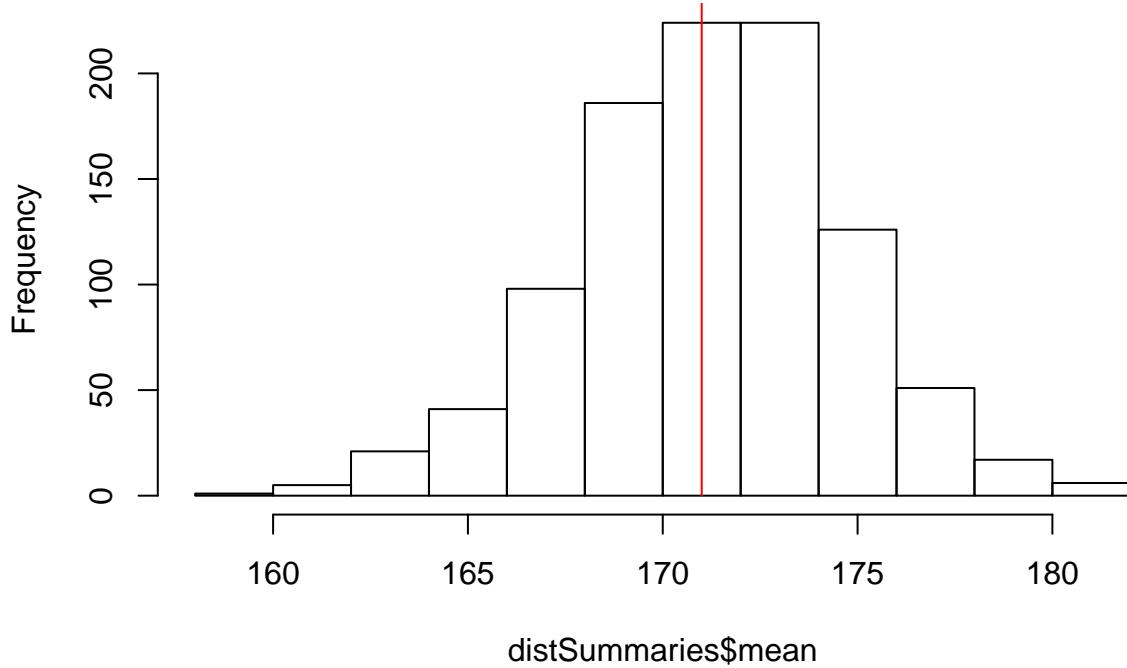
  # record summaries of distribution of Y to compare to observed Y
  distSummaries[i,]$mean = mean(ppcY)
  distSummaries[i,]$variance = var(ppcY)

}

# examine how observed sample mean compares to values generated using simulation:
hist(distSummaries$mean, main = "PPC of Sample Mean")
lines(x = c(mean(DietData$WeightLB), mean(DietData$WeightLB)), y = c(0,250), col = 2)

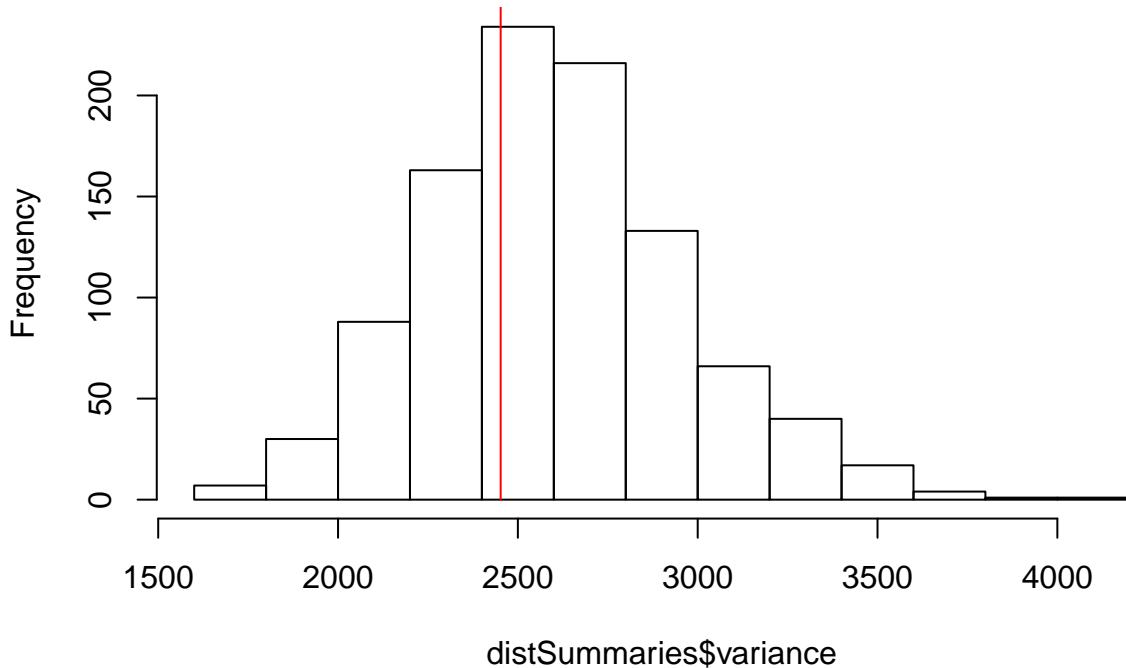
```

PPC of Sample Mean



```
# examine how observed sample mean compares to values generated using simulation:
hist(distSummaries$variance)
lines(x = c(var(DietData$WeightLB), var(DietData$WeightLB)), y = c(0,250), col = 2)
```

Histogram of distSummaries\$variance



Results

Our chains look converged, our model fit looks good, now we can look at the results:

```
# the summary function doesn't print things out well:
summary(BayesianFullModel.Posterior)
```

```
##
## Iterations = 5001:10000
## Thinning interval = 1
## Number of chains = 4
## Sample size per chain = 5000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##                                Mean        SD  Naive SE Time-series SE
## Rsquare                 9.671e-01  0.009067 6.412e-05   7.474e-05
## beta.0                  1.708e+02 48.385970 3.421e-01   3.472e-01
## beta.group2            -1.732e+02 67.877682 4.800e-01   4.723e-01
## beta.group3            -1.333e+02 64.021723 4.527e-01   4.471e-01
## beta.height             -3.862e-01  0.757297 5.355e-03   5.431e-03
## beta.height.group2     2.482e+00  1.062998 7.517e-03   7.469e-03
## beta.height.group3     3.577e+00  1.012075 7.156e-03   6.983e-03
## deviance                2.240e+02  7.032384 4.973e-02   6.654e-02
```

```

## dif.group2.group3 3.987e+01 64.521135 4.562e-01      4.605e-01
## sigma.e          1.301e+01 1.932803 1.367e-02      1.679e-02
## sigma2.e         1.730e+02 53.436527 3.779e-01      4.640e-01
## tau.e            6.285e-03 0.001783 1.261e-05      1.512e-05
##
## 2. Quantiles for each variable:
##
##              2.5%       25%       50%       75%     97.5%
## Rsquare        9.436e-01   0.96342   0.96937  9.733e-01   0.97735
## beta.0         7.393e+01 139.01303 171.02777 2.021e+02 265.44476
## beta.group2    -3.071e+02 -217.85516 -173.00514 -1.285e+02 -40.27762
## beta.group3    -2.596e+02 -175.50529 -132.88512 -9.171e+01 -5.40781
## beta.height    -1.871e+00 -0.87749  -0.38744  1.079e-01  1.12032
## beta.height.group2 3.910e-01  1.78268  2.48167  3.184e+00  4.58555
## beta.height.group3 1.562e+00  2.91743  3.57406  4.242e+00  5.57869
## deviance       2.123e+02 218.93856 223.38443 2.284e+02 239.68384
## dif.group2.group3 -8.691e+01 -2.51107  39.31963 8.169e+01 168.75507
## sigma.e         9.873e+00 11.62916 12.79293 1.413e+01 17.43055
## sigma2.e        9.748e+01 135.23742 163.65905 1.996e+02 303.82395
## tau.e           3.291e-03  0.00501  0.00611  7.394e-03  0.01026

# inspecting Bayesian Results
BayesianFullModel.Summary = summary(BayesianFullModel.Posterior)
round(BayesianFullModel.Summary$statistics, 3)

```

	Mean	SD	Naive SE	Time-series SE
## Rsquare	0.967	0.009	0.000	0.000
## beta.0	170.771	48.386	0.342	0.347
## beta.group2	-173.160	67.878	0.480	0.472
## beta.group3	-133.292	64.022	0.453	0.447
## beta.height	-0.386	0.757	0.005	0.005
## beta.height.group2	2.482	1.063	0.008	0.007
## beta.height.group3	3.577	1.012	0.007	0.007
## deviance	224.017	7.032	0.050	0.067
## dif.group2.group3	39.868	64.521	0.456	0.461
## sigma.e	13.011	1.933	0.014	0.017
## sigma2.e	173.025	53.437	0.378	0.464
## tau.e	0.006	0.002	0.000	0.000

Now, let's compare Bayesian with least squares results:

```
summary(FullModel)
```

```

##
## Call:
## lm(formula = WeightLB ~ HeightIN + factor(DietGroup) + HeightIN:factor(DietGroup),
##      data = DietData)
##
## Residuals:
##      Min      1Q      Median      3Q      Max
## -24.6724 -2.7169  0.4958  1.7918 16.4581
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 170.1664    29.2786   5.812 5.42e-06 ***
## HeightIN    -0.3768     0.4587  -0.822 0.419392

```

```

## factor(DietGroup)2      -172.5639   41.4062  -4.168 0.000345 ***
## factor(DietGroup)3      -132.6675   38.7846  -3.421 0.002241 **
## HeightIN:factor(DietGroup)2    2.4727    0.6486   3.812 0.000846 ***
## HeightIN:factor(DietGroup)3    3.5666    0.6132   5.817 5.36e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 7.949 on 24 degrees of freedom
## Multiple R-squared:  0.9787, Adjusted R-squared:  0.9742
## F-statistic: 220.3 on 5 and 24 DF,  p-value: < 2.2e-16

```

Often, each parameter is summarized by its HPD (Highest Posterior Density), a value like a confidence interval. To make our life easy, we can use the `HDIInterval` package which makes working with the R2jags output easy:

```
hdi(BayesianFullModel.Posterior)
```

```

##          Rsquare     beta.0 beta.group2 beta.group3 beta.height
## lower  0.9497963  72.73014   -306.7909 -260.431515  -1.883774
## upper  0.9787577 263.68190    -40.0764   -6.270331   1.099239
##          beta.height.group2 beta.height.group3 deviance dif.group2.group3
## lower      0.3770387           1.569665 211.2679      -92.06455
## upper      4.5577062           5.584401 237.9753      162.32140
##          sigma.e sigma2.e      tau.e
## lower   9.576525  85.14509 0.003041983
## upper  16.892455 276.51649 0.009845272
## attr(,"credMass")
## [1] 0.95

```

Reasons for differences: σ_e^2 . Note that the values reported above are the mean of σ_e^2 , which will be pulled toward the tail of the distribution. We can investigate the mode of σ_e^2 , which is more analogous to the MLE (which we get in OLS). Note, the packages below take a while to install. I have commented them out for this example, but if you would like to follow along, please uncomment and provide the time for installation.

```

#
# if (!requireNamespace("BiocManager", quietly = TRUE))
#   install.packages("BiocManager")
# BiocManager::install("genefilter", version = "3.8")
# if (!requireNamespace("modeest")) install.packages("modeest")
# library(modeest)

# mlv(x = BayesianFullModel.Posterior.All[,10], method = "naive")
# mean(BayesianFullModel.Posterior.All[,10])
#
# summary(FullModel)$sigma

```

Look at what happens if we fix the residual variance to that it was in the least squares analysis (now using the R2jags package to save some space)

```

BayesianFullModel.FixedVar = function(){
  # prior distributions:
  # Note that terms on the left are model parameter names we create here
  beta.0            ~ dnorm(betaMean.0, betaTau.0) # prior for the intercept
  beta.height       ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional slope of height
  beta.group2       ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between
  beta.group3       ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between
  beta.height.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 2
}

```

```

beta.height.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 3 (


# conditional distribution of the data (uses priors above)
for (i in 1:N){

  # creating conditional mean to put into model statement
  weight.hat[i] <- beta.0 + beta.height*height[i] + beta.group2*group2[i] + beta.group3*group3[i] + b

  # error values (for R^2)
  error[i] = weight[i]-weight.hat[i]

  # likelihood from model:
  weight[i] ~ dnorm(weight.hat[i], 1/sampleVar)

}

# calculate mean difference between groups 2 and 3:
dif.group2.group3 <- beta.group3 - beta.group2

# calculate R^2 based on http://www.stat.columbia.edu/~gelman/research/unpublished/bayes_R2.pdf
variance.pred = sd(weight.hat[])*sd(weight.hat[])
variance.error = sd(error[])*sd(error[])
Rsquare = variance.pred/(variance.pred + variance.error)
}

# the parameter names come from the model syntax in the previous step
BayesianFullModel.FixedVar.Parameters = c("beta.0",
                                         "beta.height",
                                         "beta.group2",
                                         "beta.group3",
                                         "beta.height.group2",
                                         "beta.height.group3",
                                         "dif.group2.group3",
                                         "Rsquare")

# add data and N to list object that we will pass to JAGS
BayesianFullModel.FixedVar.Data = list(N = nrow(DietData),
                                       weight = BayesianFullModel.Matrix$weight,
                                       height = BayesianFullModel.Matrix$height,
                                       group2 = BayesianFullModel.Matrix$group2,
                                       group3 = BayesianFullModel.Matrix$group3,
                                       betaMean.0 = betaMean.0,
                                       betaTau.0 = betaTau.0,
                                       sampleVar = summary(FullModel)$sigma^2)

# pick a number: Here is the date I created this syntax
# in order to have reproducible Markov chains, we need to initialize the random number generator and set
# pick a number: Here is the date I created this syntax
BayesianFullModel.FixedVar.Seed = 11022019

# Note: here, the random number seed cannot be the same per seed or the chains will be the same
RNGname = c("Wichmann-Hill", "Marsaglia-Multicarry", "Super-Duper", "Mersenne-Twister")

```

```

if (RNGname[1] %in% c("Wichmann-Hill", "Marsaglia-Multicarry",
  "Super-Duper", "Mersenne-Twister")) {
  RNGname[1] <- paste("base::", RNGname[1], sep = "")
}

BayesianFullModel.FixedVar.Init.Values <- vector("list", n.chains)
for (i in 1:n.chains) {
  BayesianFullModel.FixedVar.Init.Values[[i]]$.RNG.name <- RNGname[1]
  BayesianFullModel.FixedVar.Init.Values[[i]]$.RNG.seed <- BayesianFullModel.FixedVar.Seed + i
}

BayesianFullModelFixedVar.JAGS = jags(data = BayesianFullModel.FixedVar.Data,
  inits = BayesianFullModel.FixedVar.Init.Values,
  parameters.to.save = BayesianFullModel.FixedVar.Parameters,
  model.file = BayesianFullModel.FixedVar,
  n.chains = 4,
  n.iter = 50000,
  n.burnin = 20000,
  n.thin = 1)

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
##   Observed stochastic nodes: 30
##   Unobserved stochastic nodes: 6
##   Total graph size: 256
##
## Initializing model
BayesianFullModelFixedVar.JAGS

## Inference for Bugs model at "/var/folders/sl/n1qxb2m57pqfs1hcfqb8p4nx6fm6zn/T//Rtmpprovcrx/modeld44261"
## 4 chains, each with 50000 iterations (first 20000 discarded)
## n.sims = 120000 iterations saved
##          mu.vect  sd.vect    2.5%     25%     50%     75%
## Rsquare        0.974    0.003    0.966    0.973    0.975    0.977
## beta.0         170.039   29.343   112.368   150.181   170.113   189.853
## beta.group2   -172.491   41.478  -253.637  -200.511  -172.469  -144.496
## beta.group3   -132.650   38.894  -208.785  -158.904  -132.623  -106.323
## beta.height    -0.375    0.460    -1.272    -0.684    -0.376    -0.064
## beta.height.group2  2.471    0.650    1.194    2.034    2.472    2.910
## beta.height.group3  3.566    0.615    2.364    3.149    3.566    3.981
## dif.group2.group3  39.841   38.704  -35.723   13.656   39.862   65.994
## deviance       209.537   3.473   204.758   206.979   208.888   211.381
##          97.5%  Rhat  n.eff
## Rsquare        0.978  1.001 120000
## beta.0         227.376  1.001 110000
## beta.group2   -91.228  1.001 120000
## beta.group3   -56.789  1.001 120000
## beta.height     0.528  1.001 110000
## beta.height.group2  3.740  1.001 120000
## beta.height.group3  4.767  1.001 120000

```

```

## dif.group2.group3 115.918 1.001 110000
## deviance          217.994 1.001 120000
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 6.0 and DIC = 215.6
## DIC is an estimate of expected predictive error (lower deviance is better).

```

Model Comparison

To show how to compare two models using the DIC (Deviance Information Criterion), we will again revisit the full model, only this time we will model the variance using a log-linear model, making this a location-scale model:

$$\log\left(\frac{1}{\sigma_e^2}\right) = \log(1) - \log(\sigma_e^2) = -\log(\sigma_e^2) = \gamma_0$$

Here, γ_0 is a continuous variable that can take be real number. So, we can use a normal prior distribution on γ_0 . Note, γ_0 is the value for the intercept of the residual precision $\tau_e = \frac{1}{\sigma_e^2}$ whereas $-\gamma_0$ is the value for the intercept of the variance. As such, this analysis yields a log-normal prior distribution on σ^2 , so we are essentially comparing models with different priors.

```

BayesianFullModel = function(){
  # prior distributions:
  # Note that terms on the left are model parameter names we create here
  beta.0 ~ dnorm(betaMean.0, betaTau.0) # prior for the intercept
  beta.height ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional slope of height
  beta.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 2 and group 1
  beta.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 3 and group 1
  beta.height.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 2
  beta.height.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 3
  tau.e ~ dgamma(alpha.tau.0, beta.tau.0)      #prior for 1/sigma^2_e, residual precision

  # conditional distribution of the data (uses priors above)
  for (i in 1:N){

    # creating conditional mean to put into model statement
    weight.hat[i] <- beta.0 + beta.height*height[i] + beta.group2*group2[i] + beta.group3*group3[i] +
      beta.height.group2*height[i]*group2[i] + beta.height.group3*height[i]*group3[i]

    # error values (for R^2)
    error[i] = weight[i]-weight.hat[i]

    # likelihood from model:
    weight[i] ~ dnorm(weight.hat[i], tau.e)

  }

  # created values to track throughout the Markov chain
  sigma2.e <- 1/tau.e                                #create residual variance from 1/precision
  sigma.e   <- sqrt(sigma2.e)                         #create residual SD
}

```

```

# calculate mean difference between groups 2 and 3:
dif.group2.group3 <- beta.group3 - beta.group2

#difRPrec.group2.group3 <- gamma.group3 - gamma.group2

# calculate R^2 based on http://www.stat.columbia.edu/~gelman/research/unpublished/bayes_R2.pdf
variance.pred = sd(weight.hat[])*sd(weight.hat[])
variance.error = sd(error[])*sd(error[])
Rsquare = variance.pred/(variance.pred + variance.error)
}

BayesianFullModel.Results = jags(data = BayesianFullModel.JAGS.Data, inits = BayesianFullModel.Init.Values,
                                 parameters.to.save = BayesianFullModel.Parameters,
                                 model.file = BayesianFullModel, n.chains = 4, n.iter = 50000, n.burnin = 10000)

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
## Observed stochastic nodes: 30
## Unobserved stochastic nodes: 7
## Total graph size: 259
##
## Initializing model
## Warning in FUN(X[[i]], ...): Failed to set trace monitor for deviance
## Monitor already exists and cannot be duplicated

BayesianFullModel2 = function(){
  # prior distributions:
  # Note that terms on the left are model parameter names we create here
  beta.0 ~ dnorm(betaMean.0, betaTau.0) # prior for the intercept
  beta.height ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional slope of height
  beta.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 2 and group 1
  beta.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 3 and group 1
  beta.height.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 2
  beta.height.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 3
  gamma.0 ~ dnorm(betaMean.0, betaTau.0) #prior for the log residual precision
  tau.e <- exp(gamma.0)

  # conditional distribution of the data (uses priors above)
  for (i in 1:N){

    # creating conditional mean to put into model statement
    weight.hat[i] <- beta.0 + beta.height*height[i] + beta.group2*group2[i] + beta.group3*group3[i] +
      beta.height.group2*group2[i]*height[i] + beta.height.group3*group3[i]*height[i]

    # error values (for R^2)
    error[i] = weight[i]-weight.hat[i]

    # likelihood from model:
    weight[i] ~ dnorm(weight.hat[i], tau.e)
  }
}

```

```

# created values to track throughout the Markov chain
sigma2.e <- 1/tau.e                      #create residual variance from 1/precision
sigma.e   <- sqrt(sigma2.e)                #create residual SD

# calculate mean difference between groups 2 and 3:
dif.group2.group3 <- beta.group3 - beta.group2

#difRPrec.group2.group3 <- gamma.group3 - gamma.group2

# calculate R^2 based on http://www.stat.columbia.edu/~gelman/research/unpublished/bayes_R2.pdf
variance.pred = sd(weight.hat[])*sd(weight.hat[])
variance.error = sd(error[])*sd(error[])
Rsquare = variance.pred/(variance.pred + variance.error)
}

BayesianFullModel.Results2 = jags(data = BayesianFullModel.JAGS.Data, inits = BayesianFullModel.Init.Values,
                                   parameters.to.save = c(BayesianFullModel.Parameters, "gamma.0"),
                                   model.file = BayesianFullModel2, n.chains = 4, n.iter = 50000, n.burnin = 10000)

## Warning in jags.model(model.file, data = data, inits = init.values,
## n.chains = n.chains, : Unused variable "alpha.tau.0" in data

## Warning in jags.model(model.file, data = data, inits = init.values,
## n.chains = n.chains, : Unused variable "beta.tau.0" in data

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
##   Observed stochastic nodes: 30
##   Unobserved stochastic nodes: 7
##   Total graph size: 258
##
## Initializing model

## Warning in FUN(X[[i]], ...): Failed to set trace monitor for deviance
## Monitor already exists and cannot be duplicated

BayesianFullModel3 = function(){
  # prior distributions:
  # Note that terms on the left are model parameter names we create here
  beta.0          ~ dnorm(betaMean.0, betaTau.0) # prior for the intercept
  beta.height     ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional slope of height
  beta.group2     ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 2 and 1
  beta.group3     ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 3 and 1
  beta.height.group2 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 2
  beta.height.group3 ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of height with group 3
  gamma.0         ~ dnorm(betaMean.0, betaTau.0) #prior for the interaction of residual precision
  gamma.group2    ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 2 and 1
  gamma.group3    ~ dnorm(betaMean.0, betaTau.0) #prior for the conditional mean difference between group 3 and 1
  # tau.e <- exp(gamma.0)

  # conditional distribution of the data (uses priors above)
  for (i in 1:N){
}

```

```

# creating conditional mean to put into model statement
weight.hat[i] <- beta.0 + beta.height*height[i] + beta.group2*group2[i] + beta.group3*group3[i] + 

# error values (for R^2)
error[i] = weight[i]-weight.hat[i]

varhat[i] = gamma.0 + gamma.group2*group2[i] + gamma.group3*group3[i]
tau.e[i] <- exp(varhat[i])
# likelihood from model:
weight[i] ~ dnorm(weight.hat[i], tau.e[i])

}

# calculate mean difference between groups 2 and 3:
dif.group2.group3 <- beta.group3 - beta.group2

difRPrec.group2.group3 <- gamma.group3 - gamma.group2

# calculate R^2 based on http://www.stat.columbia.edu/~gelman/research/unpublished/bayes_R2.pdf
# variance.pred = sd(weight.hat[])*sd(weight.hat[])
# variance.error = sd(error[])*sd(error[])
# Rsquare = variance.pred/(variance.pred + variance.error)
}

BayesianFullModel3.Parameters = c("beta.0",
                                 "beta.height",
                                 "beta.group2",
                                 "beta.group3",
                                 "beta.height.group2",
                                 "beta.height.group3",
                                 "dif.group2.group3",
                                 "difRPrec.group2.group3",
                                 "gamma.0", "gamma.group2", "gamma.group3")

BayesianFullModel.Results3 = jags(data = BayesianFullModel.JAGS.Data, inits = BayesianFullModel.Init.Values,
                                   parameters.to.save = BayesianFullModel.FixedVar.Parameters,
                                   model.file = BayesianFullModel2, n.chains = 4, n.iter = 50000, n.burnin = 10000)

## Warning in jags.model(model.file, data = data, inits = init.values,
## n.chains = n.chains, : Unused variable "alpha.tau.0" in data

## Warning in jags.model(model.file, data = data, inits = init.values,
## n.chains = n.chains, : Unused variable "beta.tau.0" in data

## Compiling model graph
## Resolving undeclared variables
## Allocating nodes
## Graph information:
##   Observed stochastic nodes: 30
##   Unobserved stochastic nodes: 7
##   Total graph size: 258
##
## Initializing model

```

BayesianFullModel.Results

```
## Inference for Bugs model at "/var/folders/sl/n1qxb2m57pqfs1hcfqb8p4nx6fm6zn/T//RtmpprovRx/modeld4427"
## 4 chains, each with 50000 iterations (first 20000 discarded)
## n.sims = 120000 iterations saved
##          mu.vect sd.vect    2.5%    25%    50%    75%
## Rsquare      0.967  0.009   0.944   0.963   0.969   0.973
## beta.0       170.287 48.242  74.779 138.618 170.231 202.002
## beta.group2  -172.571 68.417 -307.151 -217.666 -172.642 -127.937
## beta.group3  -133.009 64.039 -259.031 -175.175 -133.203 -90.934
## beta.height   -0.379  0.756  -1.867  -0.875  -0.378  0.117
## beta.height.group2  2.473  1.072  0.348   1.771  2.475  3.178
## beta.height.group3  3.572  1.013  1.564   2.909  3.575  4.238
## dif.group2.group3  39.561 64.220 -86.651 -2.403  39.237 81.724
## sigma.e        12.985  1.926   9.873 11.620 12.761 14.100
## sigma2.e       172.313 53.282  97.471 135.024 162.842 198.797
## tau.e          0.006  0.002   0.003   0.005  0.006  0.007
## deviance       223.906 6.952  212.282 218.914 223.253 228.162
##          97.5%  Rhat  n.eff
## Rsquare        0.977 1.001 120000
## beta.0         265.233 1.001 53000
## beta.group2   -37.272 1.001 32000
## beta.group3   -6.392 1.001 120000
## beta.height    1.119 1.001 51000
## beta.height.group2  4.584 1.001 32000
## beta.height.group3  5.568 1.001 120000
## dif.group2.group3 166.045 1.001 39000
## sigma.e        17.368 1.001 92000
## sigma2.e       301.659 1.001 92000
## tau.e          0.010 1.001 92000
## deviance       239.272 1.001 97000
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 24.2 and DIC = 248.1
## DIC is an estimate of expected predictive error (lower deviance is better).
```

BayesianFullModel.Results2

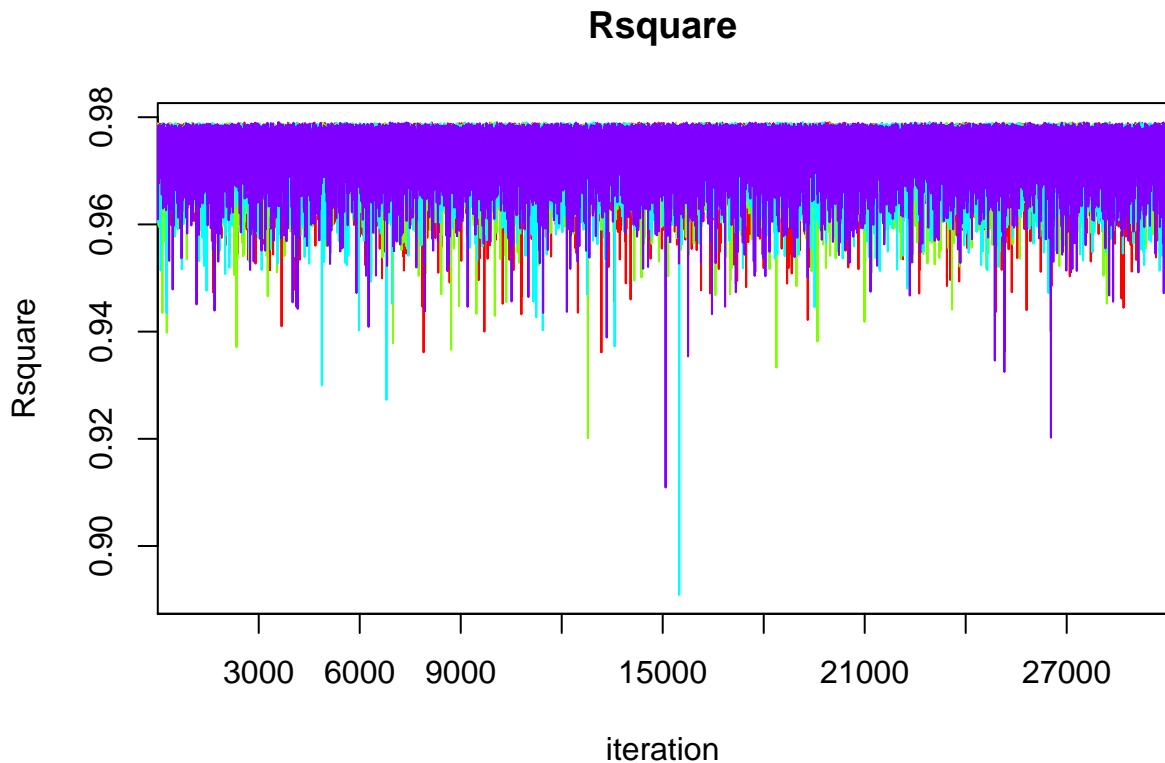
```
## Inference for Bugs model at "/var/folders/sl/n1qxb2m57pqfs1hcfqb8p4nx6fm6zn/T//RtmpprovRx/modeld4423"
## 4 chains, each with 50000 iterations (first 20000 discarded)
## n.sims = 120000 iterations saved
##          mu.vect sd.vect    2.5%    25%    50%    75%
## Rsquare      0.974  0.004   0.964   0.972   0.975   0.977
## beta.0       170.148 30.595  109.582 149.987 170.160 190.338
## beta.group2  -172.635 43.152 -257.670 -200.965 -172.775 -144.224
## beta.group3  -132.611 40.475 -212.900 -159.262 -132.524 -106.110
## beta.height   -0.377  0.479  -1.321  -0.693  -0.377  -0.061
## beta.height.group2  2.474  0.676   1.143   2.031   2.477  2.917
## beta.height.group3  3.566  0.640   2.305   3.146   3.565  3.985
## dif.group2.group3  40.024 40.522 -40.271  13.335  40.035 66.769
## gamma.0        -4.190  0.296  -4.807  -4.381  -4.176  -3.984
## sigma.e        8.214  1.248   6.206   7.330   8.069  8.941
```

```

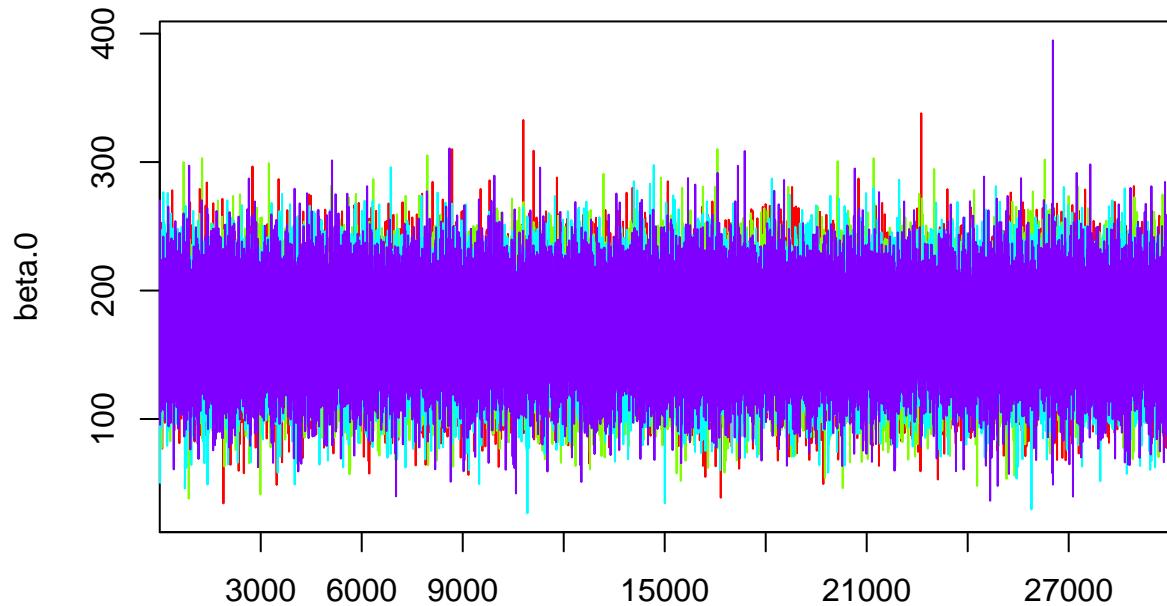
## sigma2.e          69.033 21.871 38.519 53.733 65.101 79.944
## tau.e            0.016  0.005  0.008  0.013  0.015  0.019
## deviance         210.823 4.268 204.769 207.694 210.081 213.150
##                97.5%   Rhat  n.eff
## Rsquare           0.978 1.001 99000
## beta.0            230.414 1.001 120000
## beta.group2       -87.621 1.001 120000
## beta.group3       -52.758 1.001 120000
## beta.height        0.575 1.001 120000
## beta.height.group2 3.806 1.001 120000
## beta.height.group3 4.836 1.001 120000
## dif.group2.group3 119.693 1.001 120000
## gamma.0           -3.651 1.001 120000
## sigma.e            11.060 1.001 120000
## sigma2.e           122.313 1.001 120000
## tau.e              0.026 1.001 120000
## deviance          221.141 1.001 120000
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 9.1 and DIC = 219.9
## DIC is an estimate of expected predictive error (lower deviance is better).

traceplot(BayesianFullModel.Results3)

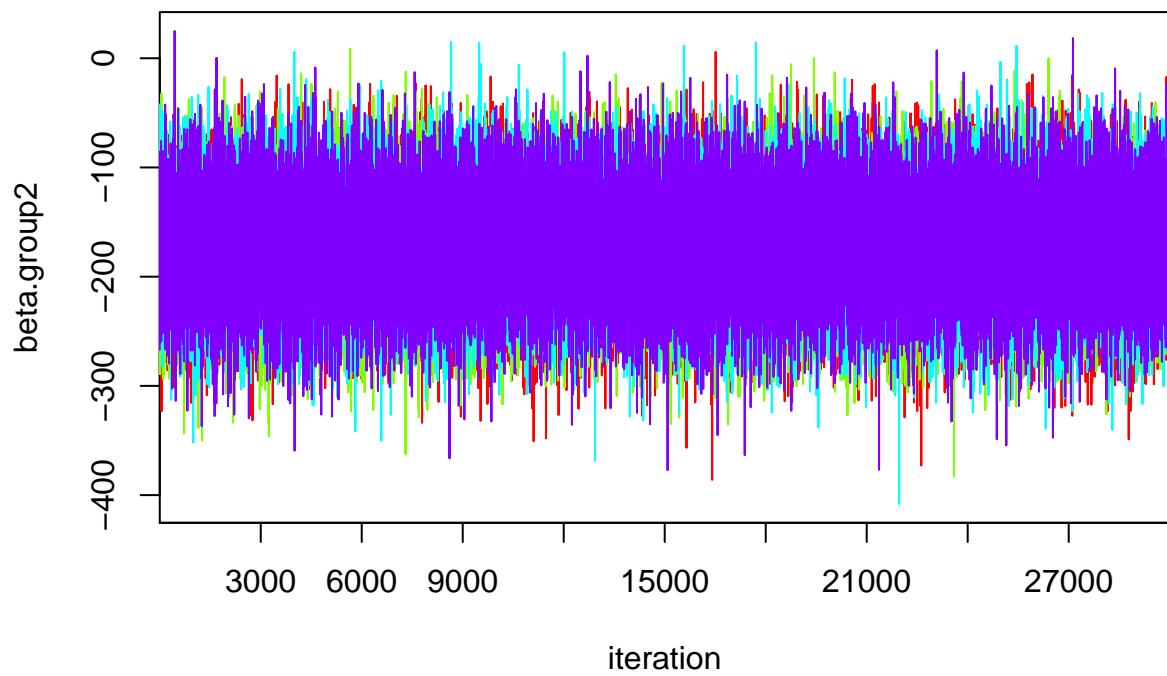
```



beta.0

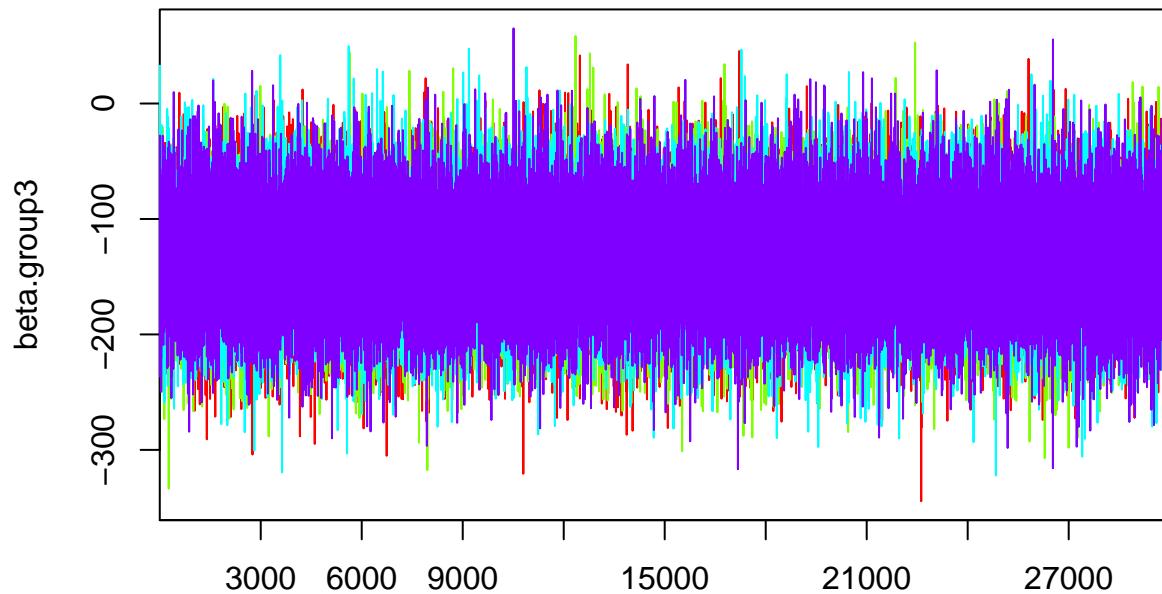


iteration
beta.group2

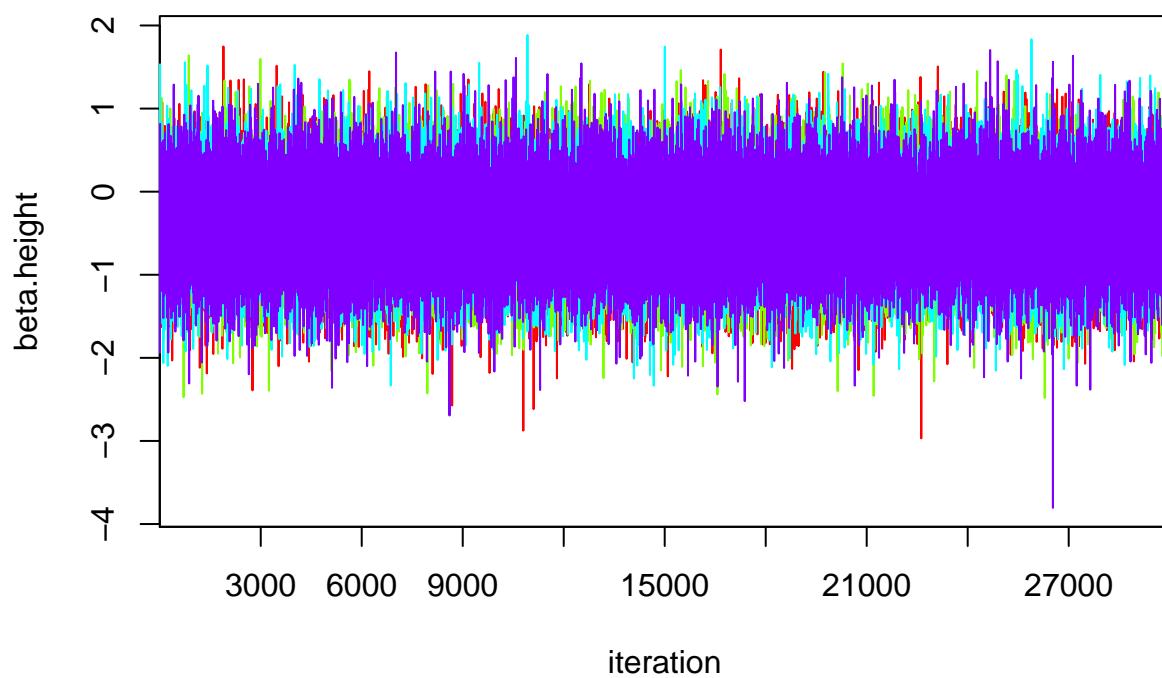


iteration

beta.group3

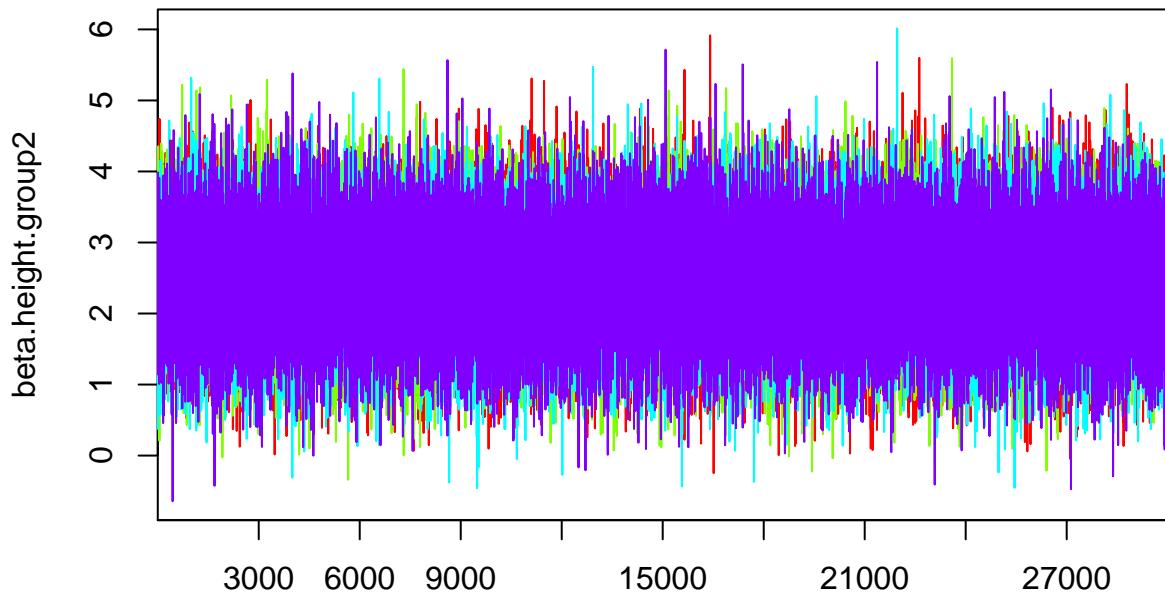


iteration
beta.height

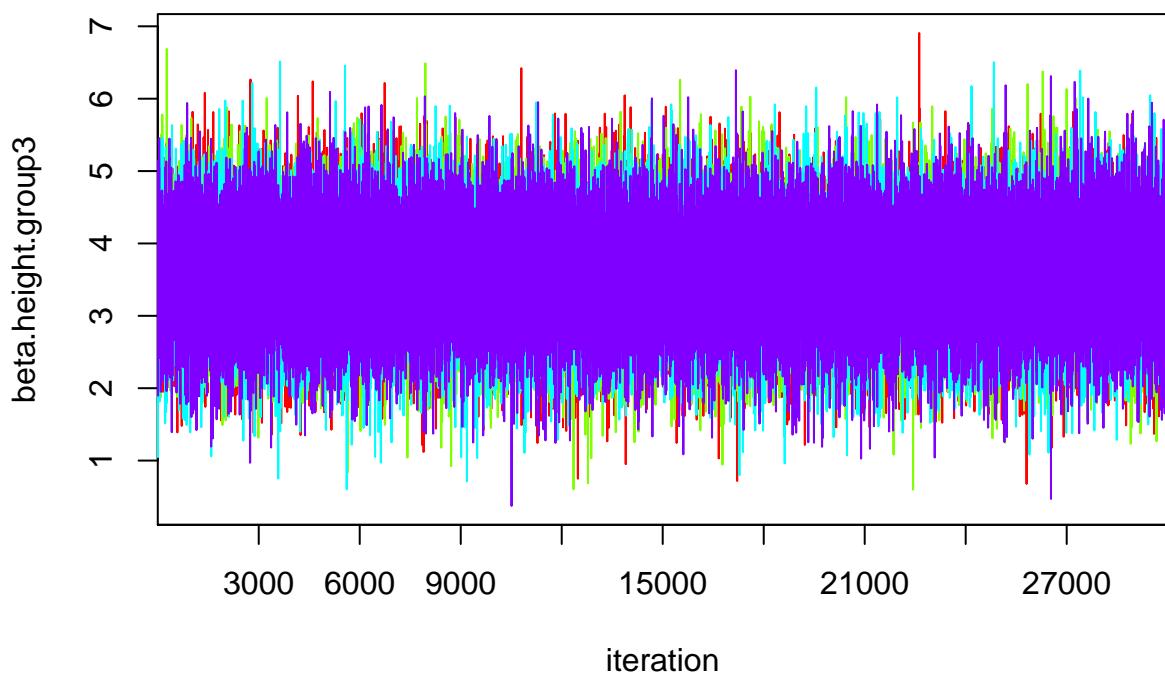


iteration

beta.height.group2

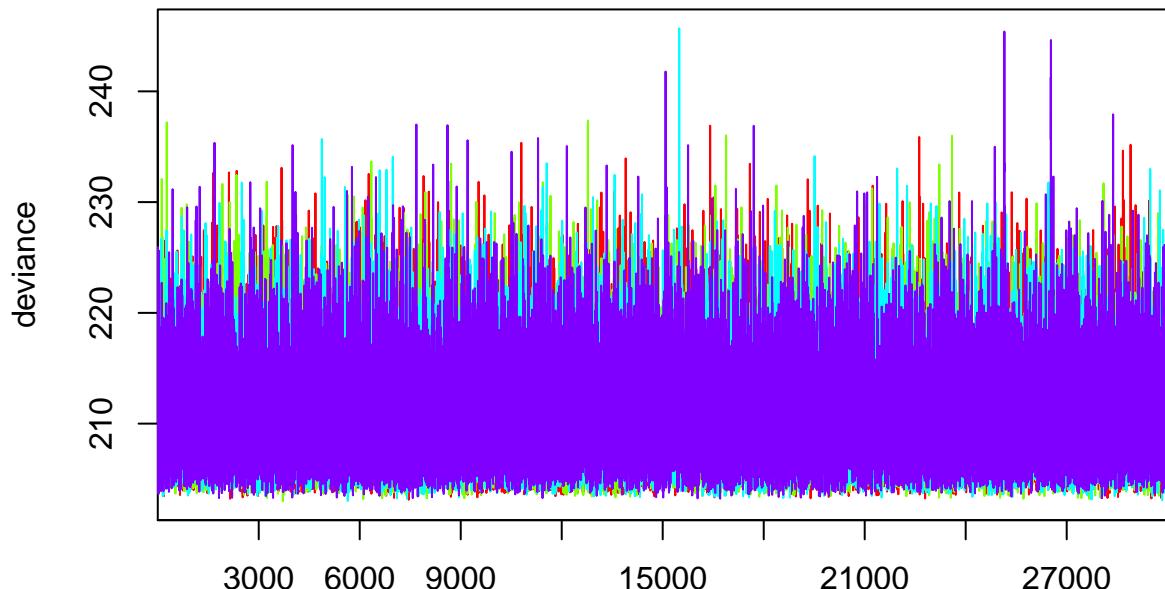


iteration
beta.height.group3

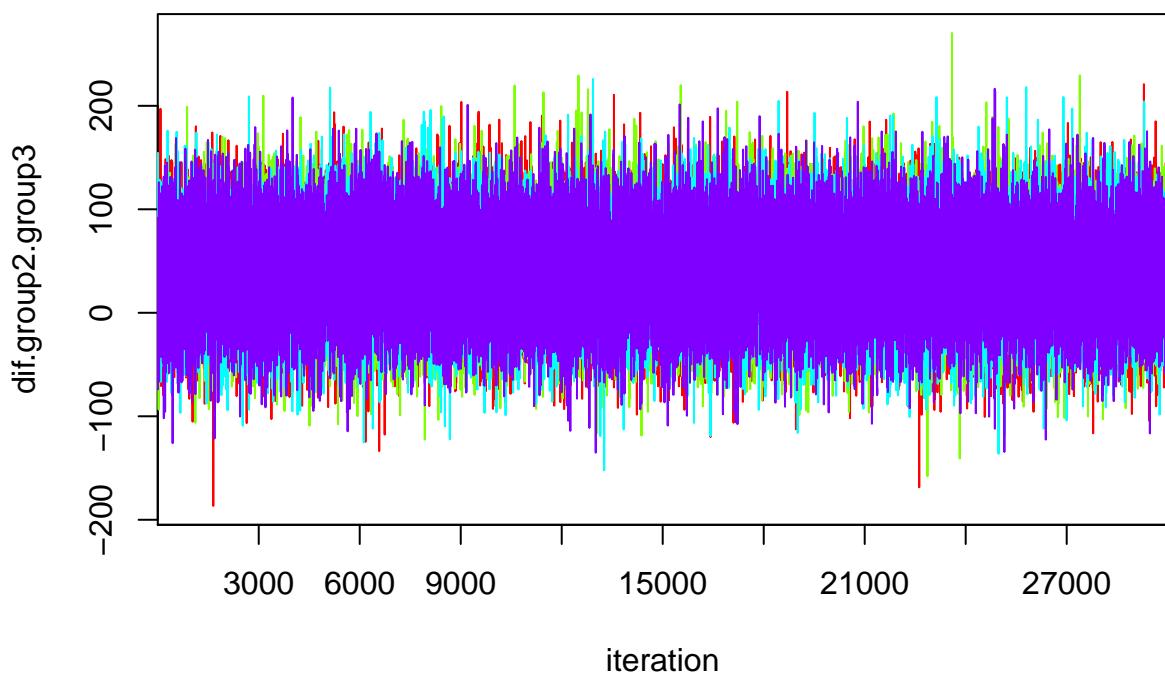


iteration

deviance



iteration
dif.group2.group3



iteration

BayesianFullModel.Results3

```
## Inference for Bugs model at "/var/folders/sl/n1qxb2m57pqfs1hcfqb8p4nx6fm6zn/T//Rtmpprovcrx/modeld4429a"
## 4 chains, each with 50000 iterations (first 20000 discarded)
## n.sims = 120000 iterations saved
##          mu.vect    sd.vect      2.5%     25%     50%     75%
##
```

```

## Rsquare          0.974   0.004   0.964   0.972   0.975   0.977
## beta.0          170.063  30.547  109.756  149.972  170.072  190.103
## beta.group2     -172.447  43.213 -257.774 -200.712 -172.401 -143.937
## beta.group3     -132.535  40.432 -212.547 -158.961 -132.575 -106.017
## beta.height      -0.375   0.479   -1.321   -0.690   -0.376   -0.061
## beta.height.group2  2.471   0.677   1.126   2.023   2.471   2.914
## beta.height.group3  3.565   0.639   2.298   3.145   3.566   3.982
## dif.group2.group3 39.912   40.531  -40.295  13.241  39.931  66.571
## deviance         210.783  4.260   204.739  207.665  210.019  213.103
##
##               97.5% Rhat n.eff
## Rsquare          0.978 1.001 93000
## beta.0           230.395 1.001 120000
## beta.group2     -86.942 1.001 55000
## beta.group3     -52.581 1.001 100000
## beta.height      0.569 1.001 120000
## beta.height.group2 3.810 1.001 58000
## beta.height.group3 4.825 1.001 79000
## dif.group2.group3 120.023 1.001 120000
## deviance        221.030 1.001 45000
##
## For each parameter, n.eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor (at convergence, Rhat=1).
##
## DIC info (using the rule, pD = var(deviance)/2)
## pD = 9.1 and DIC = 219.9
## DIC is an estimate of expected predictive error (lower deviance is better).

```

Algorithm Used

Because we changed away from a conjugate prior, we no longer use Gibbs sampling for all parameters. Instead, we now have a “RealSlicer” (a form of Metropolis-Hastings) for γ_0 :

```

list.samplers(BayesianFullModel.Results$model)

## $`glm::Generic`
## [1] "beta.height.group3" "beta.height.group2" "beta.group3"
## [4] "beta.group2"         "beta.height"       "beta.0"
##
## $`bugs::ConjugateGamma`
## [1] "tau.e"

list.samplers(BayesianFullModel.Results2$model)

## $`glm::Generic`
## [1] "beta.height.group3" "beta.height.group2" "beta.group3"
## [4] "beta.group2"         "beta.height"       "beta.0"
##
## $`base::RealSlicer`
## [1] "gamma.0"

```

Deviance Information Criterion

The DIC is a Bayesian version of an information criterion used to compare the relative fit of multiple models (see https://en.wikipedia.org/wiki/Deviance_information_criterion). Lower values are better.

The deviance is given by -2 times the likelihood plus a constant (that goes away in model comparisons).

$$D = -2 \log(p(y | \theta)) + C$$

The DIC is given by the deviance, evaluated at the EAP estimate of all parameters, plus a penalty. The earliest form of the penalty p_D is:

$$p_D = \bar{D} - D(\bar{\theta})$$

$$DIC = D(\bar{\theta}) + 2p_D$$

Here, the smallest value of the DIC belongs to the model with the log-linear variance prior. So, it is preferred and would be the one interpreted.